

# 17th edition SQLDay Conference

12-14 May 2025, WROCŁAW + ONLINE



---

## Platinum sponsors

---



---

## Gold sponsors

---



---

## Silver sponsors

---



# **Mastering Microsoft Fabric Data Warehouse Performance**

**Filip Popović**

Senior Program Manager, Microsoft, Serbia

# Agenda

- Key ingredients for optimal performance
- Query activity and insights
- Speed up your queries with data clustering
- Consumption considerations
- Ingestion best practices

# **Key ingredients for optimal performance**

# Always optimal query performance

\*Guardrails protect from overconsumption

- Fabric bursting automatically allocates resources as needed to achieve optimal performance
- Query performance is always the same regardless of SKU\*
- F2 capacity can run 1TB workloads thanks to bursting!



F64


# No more sizing for peaks

\*Guardrails protect from overconsumption



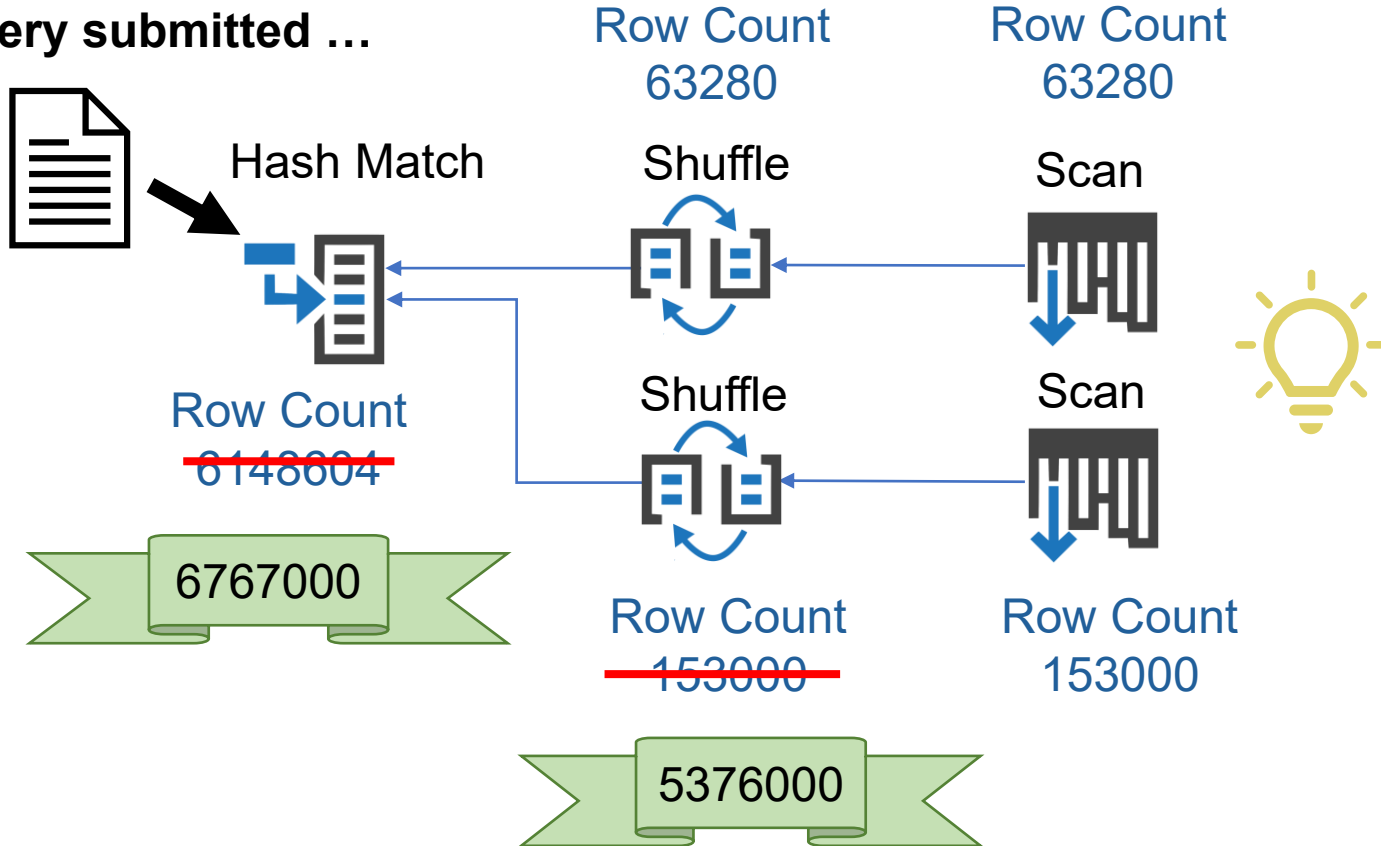
# No more sizing for peaks

- Smoothing to future time windows to avoid throttling
- Temporary peaks are handled without DBA intervention or increasing capacity size
- Use Capacity metrics app to track consumption
- F64 (trial) capacity can run TPCB 10TB power run thanks to bursting and smoothing!



# Learned resource estimation


Query submitted ...



1. Generate plan with current statistics.
2. Check if actual cardinality for sub-trees exists from previous query execution.
3. Use actual cardinality when applicable.
4. Execute query with accurate estimations!



# Learned resource estimation



Optimize resource estimates for each task in your query based on previous executions



Improve query performance



Improve compute utilization by avoiding resource miss-estimates



Reduce data spilling significantly

# Query activity and insights

# Workload monitoring

- DMVs
  - Active connections, sessions and requests
- Query Insights
  - Historical – track performance over time
  - Aggregated – more actionable insights
    - Queries are aggregated by query text excluding predicates
  - Cross-database queries are shown in context DB query insights
- Query activity
  - Monitoring UI built on top of DMVs and Query Insights

# Identify optimization opportunities

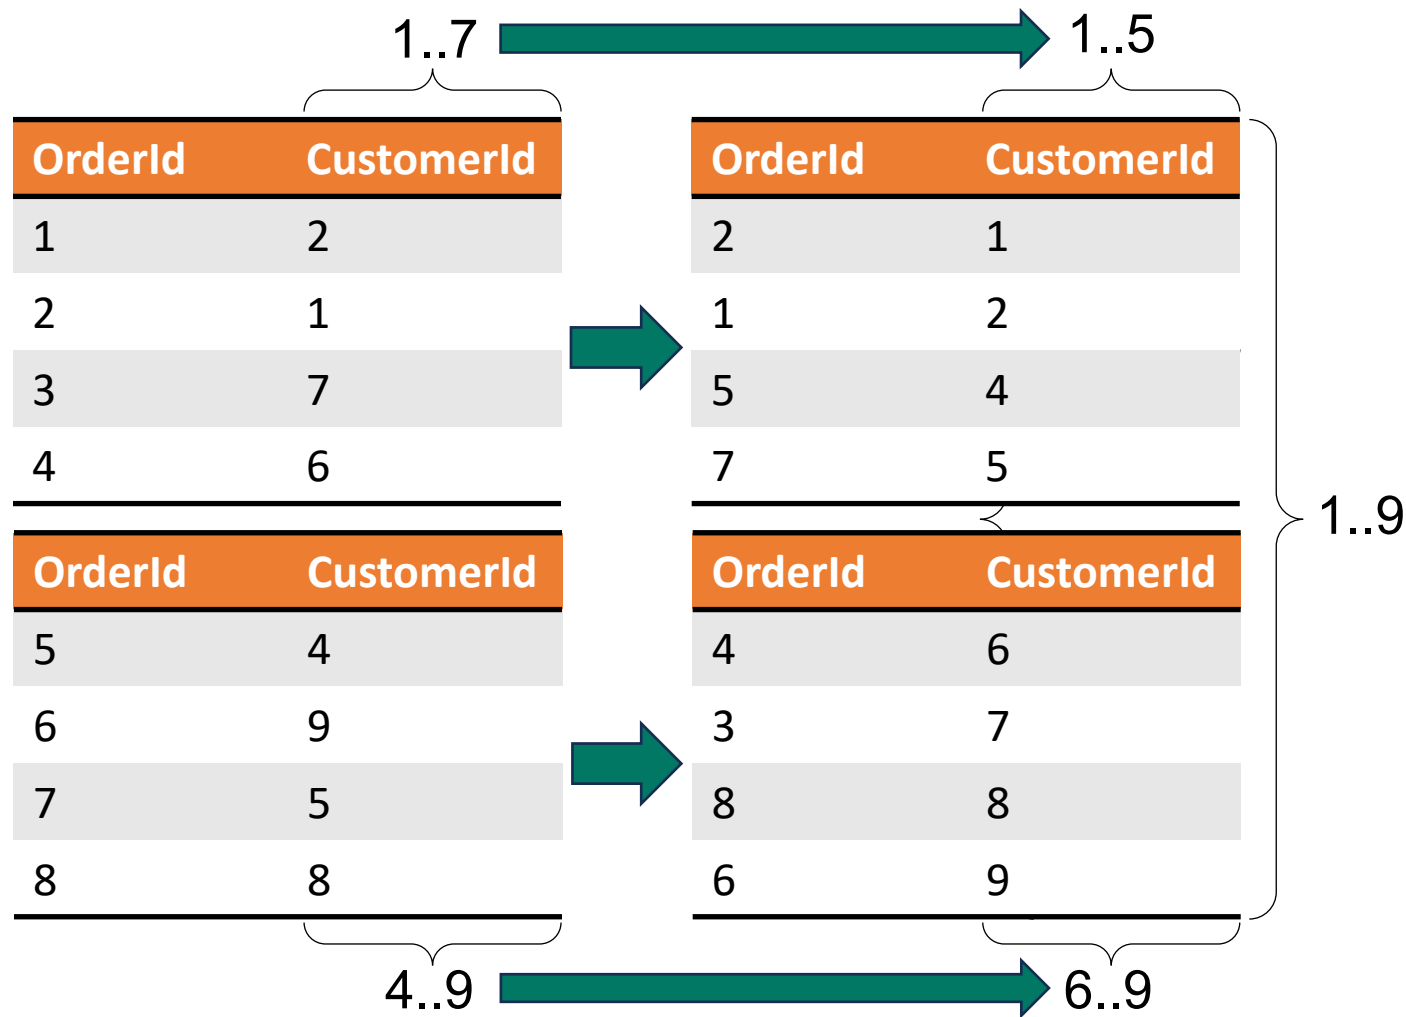
- queryinsights.exec\_requests\_history
  - data\_scanned\_remote\_storage\_mb
  - data\_scanned\_disk\_mb
  - data\_scanned\_memory\_mb
  - allocated\_cpu\_time\_ms

	label	data_scanned_remote_storage_mb	data_scanned_memory_mb	data_scanned_disk_mb	total_cpu_time_ms
1	baseline - cold	63399	0	0	17985154
2	baseline - warm	0	63399	0	3562064
3	clustering - cold	8788	0	0	1312427
4	clustering - warm	0	8788	0	270994

**Speed up your queries with data  
clustering**

**Coming soon**

# Reduce IO with data clustering



Collocate data points that are close to each other

Improve performance by skipping rowgroups and files

## Data clustering – guidelines



- Pick columns used in query predicates
- Pick column with medium to high ratio of unique values ( $\geq 60\%$ )

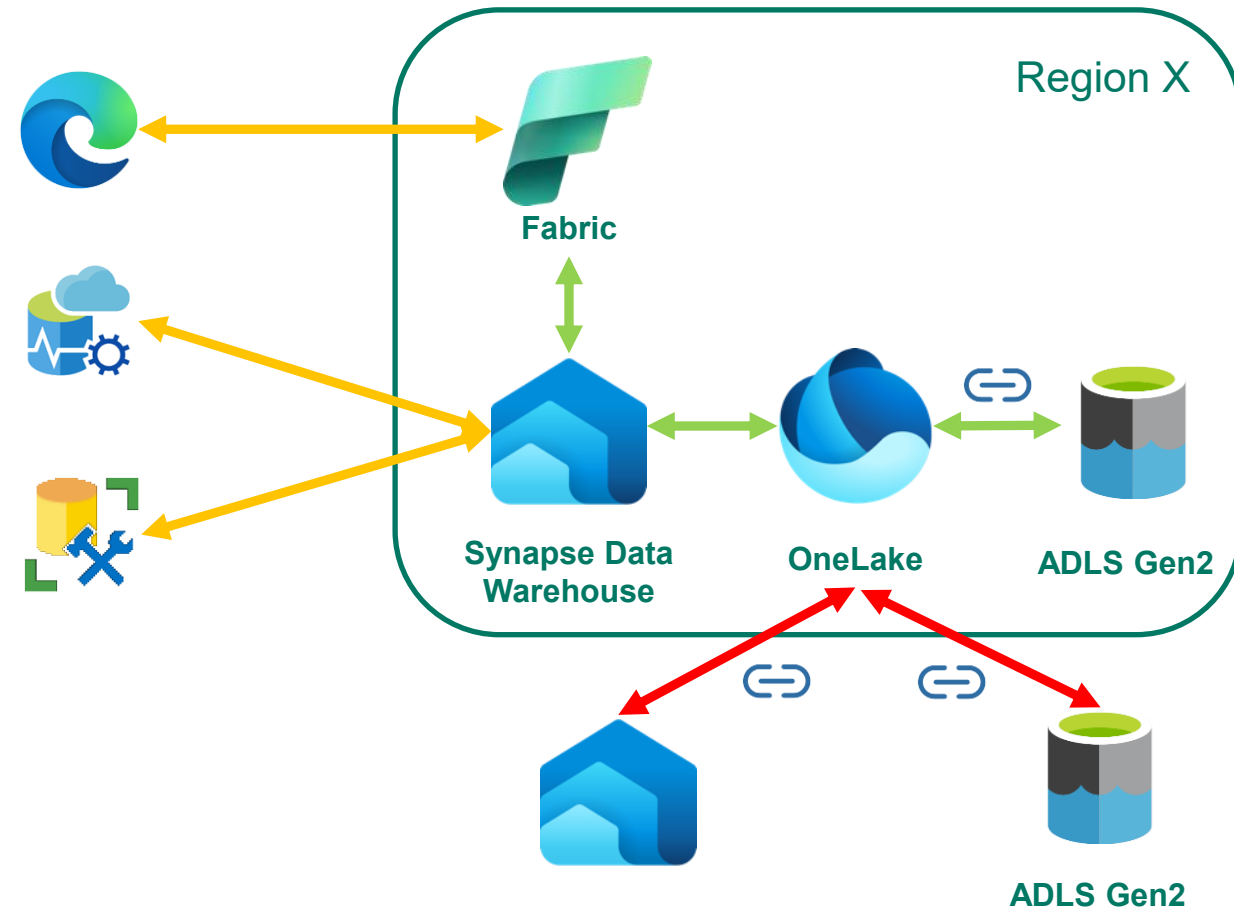
```
SELECT COUNT(DISTINCT <columnName>) * 1. / COUNT(*) FROM <table>
```

# Consumption considerations



# Where are my compute and data?

- Collocate resources
- Compute is where capacity is
- Mind network latency between:
  - The client and the endpoint
  - The engine and the data in case of shortcuts




# Choosing Optimal Data Types for Real-World Problems\*

- Use smallest data type that accommodates values
- Use proper types
- Use varchar instead of char
- Use NOT NULL instead of NULL
- Check table schema created by tools (e.g. Pipelines)

<code>year bigint</code>	<code>-&gt;</code>	<code>year int</code>
<code>amount decimal</code>	<code>-&gt;</code>	<code>amount decimal(p,s)</code>
<code>code_id varchar(10)</code>	<code>-&gt;</code>	<code>code_id int</code>
<code>name char(50)</code>	<code>-&gt;</code>	<code>name varchar(50)</code>
<code>id bigint</code>	<code>-&gt;</code>	<code>id bigint NOT NULL</code>
<code>name varchar(8000)</code>	<code>-&gt;</code>	<code>name varchar(50)</code>

# Optimize Lakehouse tables

Fabric Warehouse can only optimize Warehouse tables



Yet another table to SQL engine, same best practices apply

Check data types, particularly string lengths, decimal precision and scale and use NOT NULL if possible

OPTIMIZE table to get optimal number and sizes of files

Partition table by columns commonly used for filtering

Z-ORDER table by non-partitioning columns commonly used for filtering.

# Data ingestion best practices


# Don't ingest, shortcut or mirror instead

## Shortcuts into the lakehouse

- Makes no additional copy of the data

## Mirror data

- Near real time replication with zero ETL
- Optimal format for analytics by default
- Currently supported sources:
  - Snowflake
  - Azure Cosmos DB
  - Azure SQL DB
  - Azure SQL MI
  - Azure Database for PostgreSQL Flexible Server
  - Extensible with Open Mirroring!



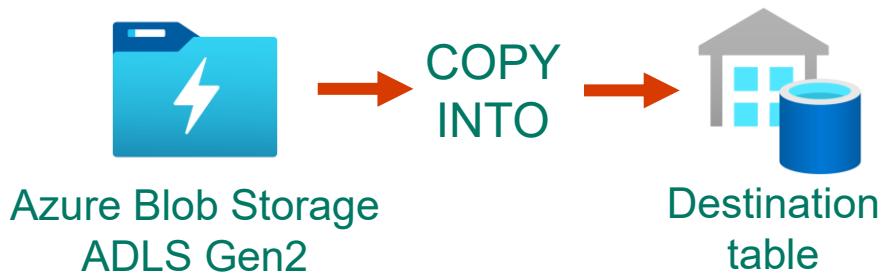
## Few tips for a start

From	Format	T-SQL statement	Capabilities
Lakehouse table	Delta	CTAS/INSERT with SELECT	Ingestion with transformation.
ADLS	Delta	Shortcut + CTAS/INSERT with SELECT	Ingestion with transformation.
ADLS	CSV, Parquet	COPY INTO	Ingestion w/o transformation.
ADLS	CSV, Parquet	CTAS/INSERT+SELECT with OPENROWSET	Ingestion with transformation of subset of rows.
ADLS	CSV	BULK INSERT	Ingestion with transformation. Migrating existing code and third-party code.

**CLONE TABLE** for instantly available replica

# COPY INTO

- Enables flexible, high-throughput data ingestion from external Azure storage accounts




```
COPY INTO table_name.[(Column_list)]
FROM '<external_location>' [,...n]
WITH
(
  [FILE_TYPE = {'CSV' | 'PARQUET' '}]
  [, CREDENTIAL = (AZURE CREDENTIAL) ]
  [, ERRORFILE =
' [http(s)://account/container]/directory[/]] ' ]
  [, ERRORFILE_CREDENTIAL = (AZURE CREDENTIAL) ]
  [, MAXERRORS = max_errors ]
  [, COMPRESSION = { 'Gzip' | 'Snappy' } ]
  [, FIELDQUOTE = 'string_delimiter' ]
  [, FIELDTERMINATOR = 'field_terminator' ]
  [, ROWTERMINATOR = 'row_terminator' ]
  [, FIRSTROW = first_row ]
  [, ENCODING = { 'UTF8' | 'UTF16' } ]
  [, PARSER_VERSION = { '1.0' | '2.0' } ]
)
```

# Ingestion best practices - size and number of files

## COPY INTO:

- More small files > few large files
- At least 4MB file size
- Keep the number of files to be at least 1,000
- Best throughput using 1,000 files is 7x faster than single large file




\*TPCH 1TB, Orders table



# Ingestion best practices - parallelize

## COPY INTO:

- Parallel loads achieve higher throughput

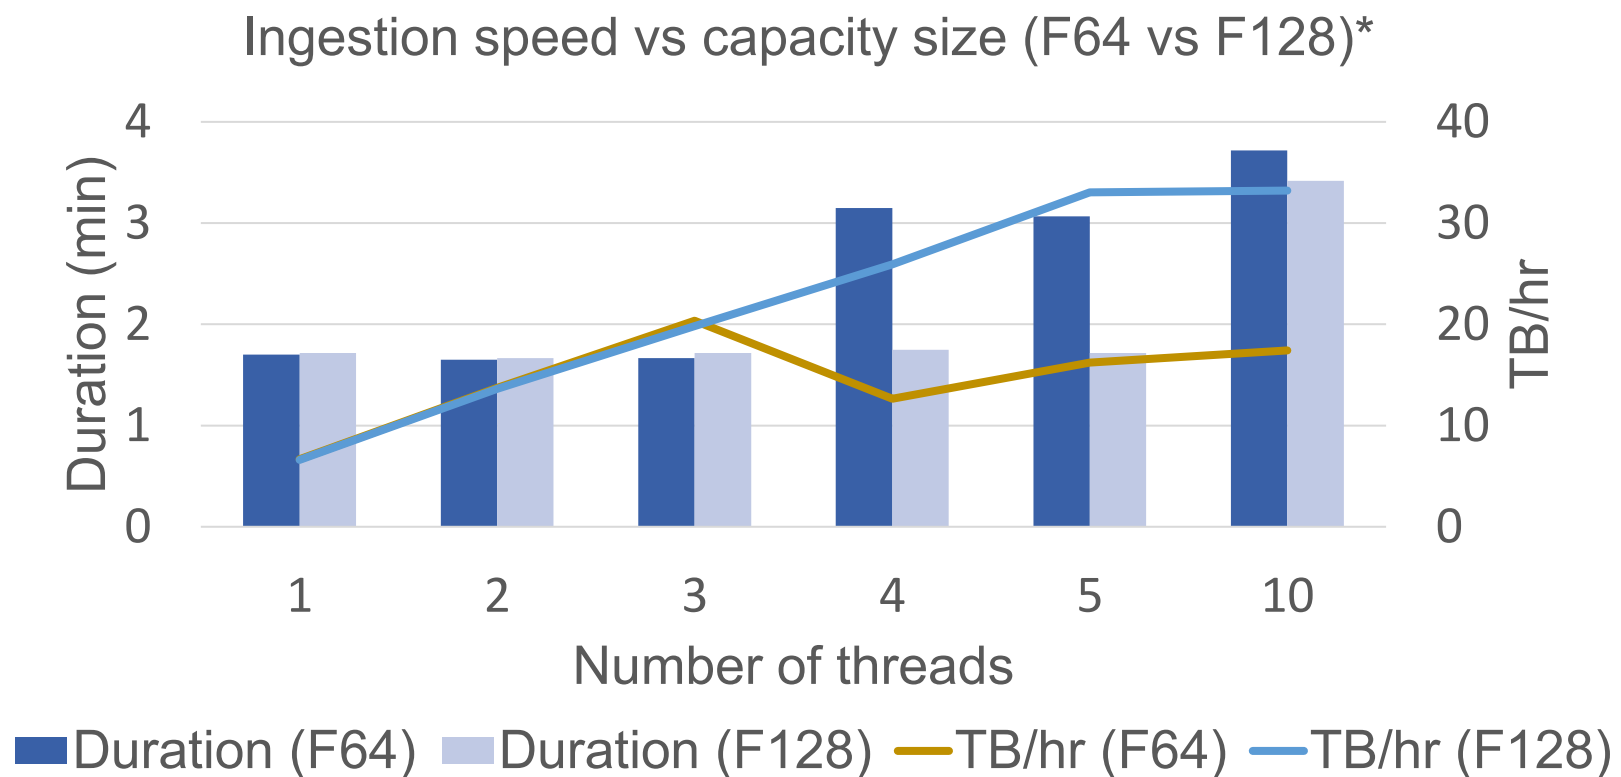


\*TPCH 1TB, Orders table, 170GB each thread

# Increase throughput with larger Fabric capacities

All ingestion (pipelines, dataflows, SQL):

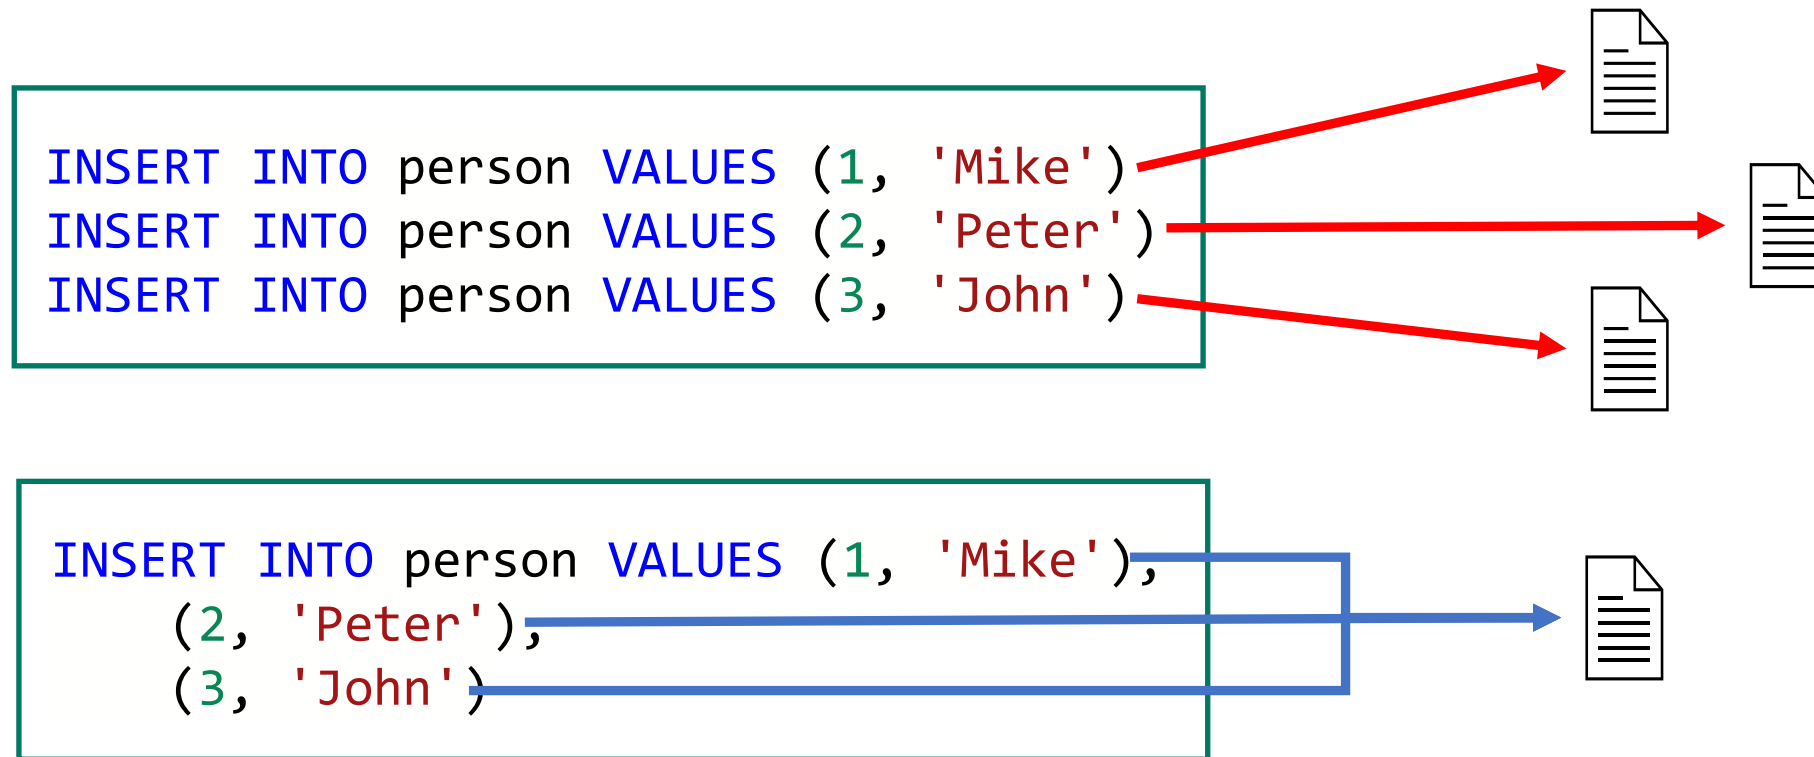
- Scales with larger Fabric Capacities



\*TPCH 1TB, Orders table, 170GB each thread

# INSERT/UPDATE/DELETE data in batches

Avoid singleton INSERT INTO statements – ineffective for both writes and reads



# Q&A, feedback and bits for refreshing memory

- Q&A
- Please provide feedback in the Whova App
- Share with your colleagues and your future self 😊
  - [Ingesting data into Fabric Warehouse](#)
  - [Performance best practices](#)

17th edition  
**SQLDay Conference**  
12-14 May 2025, WROCŁAW + ONLINE

---

Platinum sponsors

---



---

Gold sponsors

---



---

Silver sponsors

---

