

# 17th edition SQLDay Conference

12-14 May 2025, WROCŁAW + ONLINE

· Platinum sponsors





**Gold sponsors** 











**Silver sponsors** 

















# Partitioning Approaches in Database Systems:

Insights from Oracle, SQL Server, and PostgreSQL

Sarah Kossakowska, Szymon Lasota





# Szymon Lasota

### Cloud Architect @ Lingaro



https://www.linkedin.com/in/szymon-lasota/































# Sarah Kossakowska

Data Integration Engineer @ Lingaro



sarah.kossakowska@lingarogroup.com



# SQLDay 2025 Partitioning Agenda

01

Partitioning intro

02

SQL Server deep dive on Partitioning

03

Oracle deep dive on Partitioning

04

PostgreSQL deep dive on Partitioning

05

What you need to know about partitioning Azure SQL Database

06

QΑ



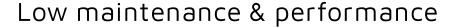


Partitioning Intro

Work only with data that you need.

VS





- Additional effort in design and planning
- More complex data loading



High maintenance & performance?

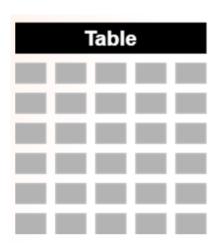
- Faster data access
- Improved data management flexibility

# Physically Divided Tables

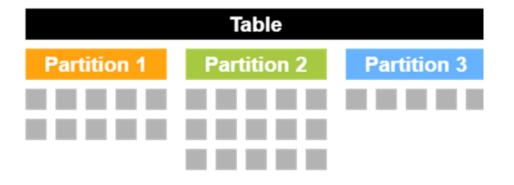


Patitioning is a physical split of large tables or indexes into smaller manageable units. Each partition can be stored, queried, and maintained separately, though they appear as single logical table to the user.

#### Non-Partitioned



#### Partitioned

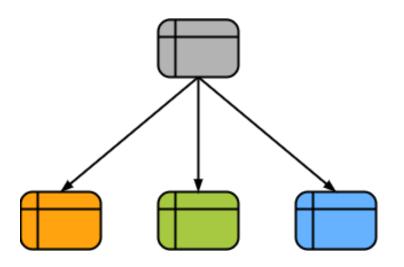


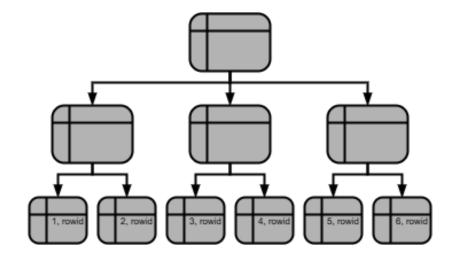


## They are not alternatives.

Size of retrieved data portions. Large volumen or specific rows.

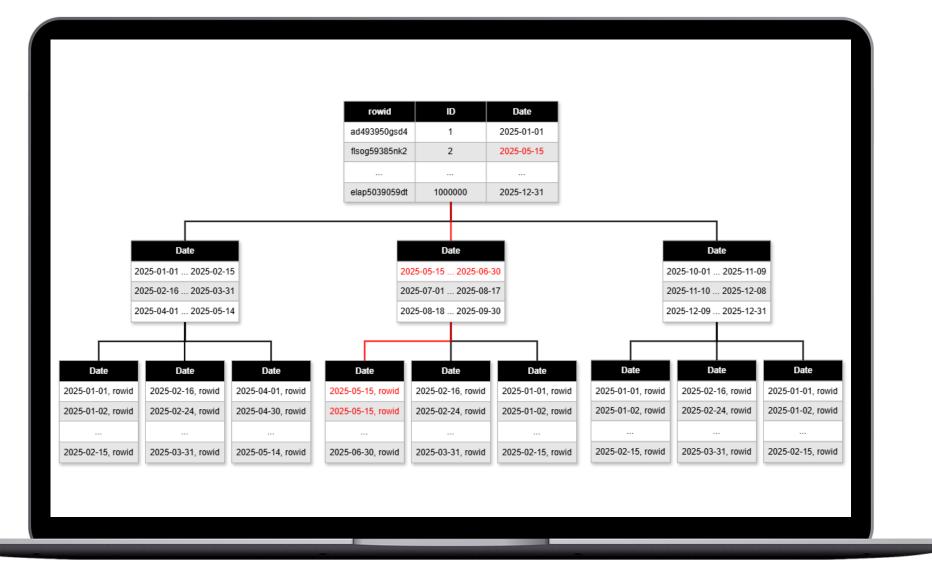
Access path. Immediately or rowid by rowid.





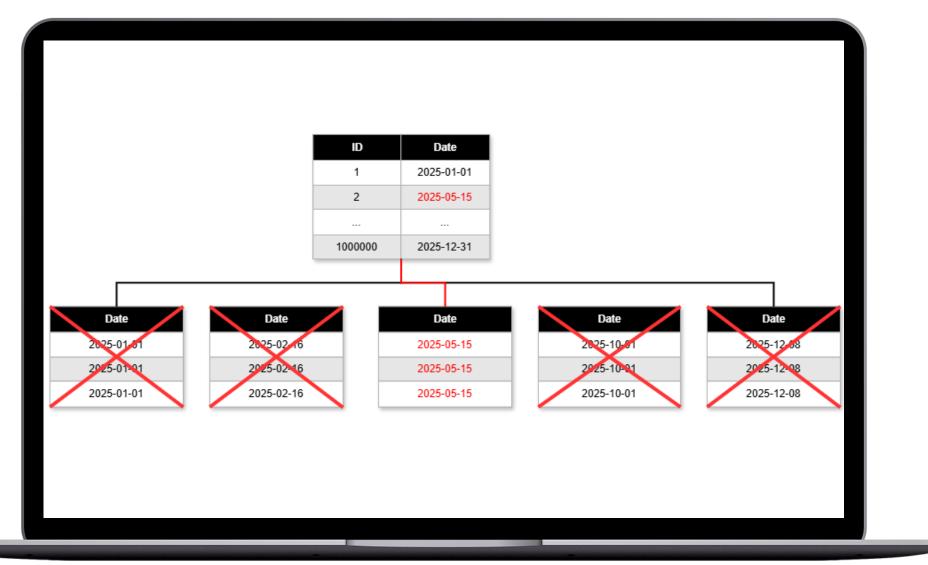










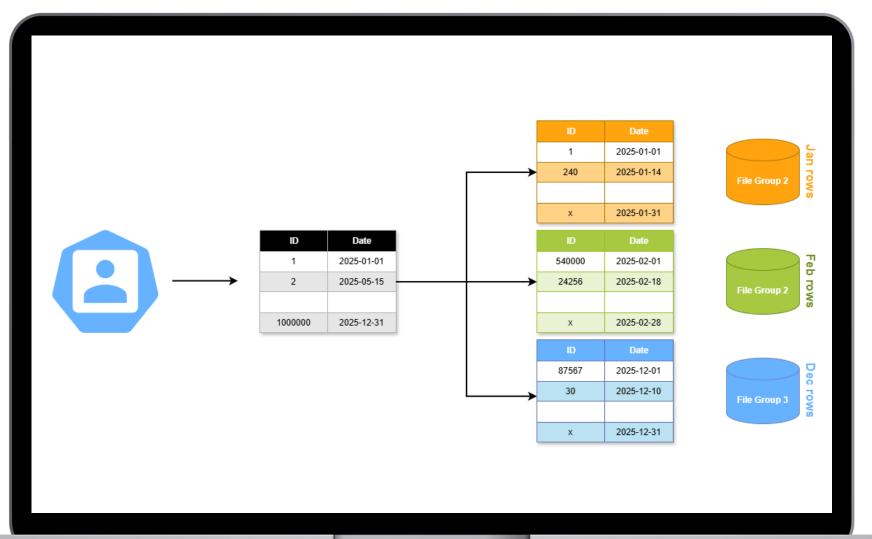


# No Impact on User View



00

Structure remains
untouched. No changes
required for user
or the application.
The table is partitioned
based on the values
of existing columns.



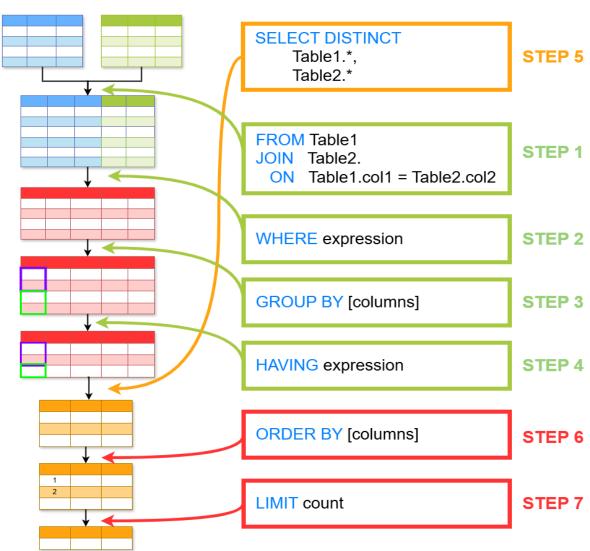
# Query Execution Order



00

Advantage of partitioning through proper query construction.

Query Planner analyzes the query's filters in the WHERE clause and JOIN conditions.



# Partition Pruning with WHERE Clause



Filter directly on the partition column using simple comparisons (e.g., =, BETWEEN, <, >) without applying functions or expressions.

- ✓ WHERE Date = '2024-01-01'
- WHERE Date BETWEEN '2024-01-01' AND '2024-12-31'
- WHERE Date >= '2024-01-01' AND order\_date <= '2024-12-31'
- WHERE EXTRACT(YEAR FROM Date) = 2024



# Demo

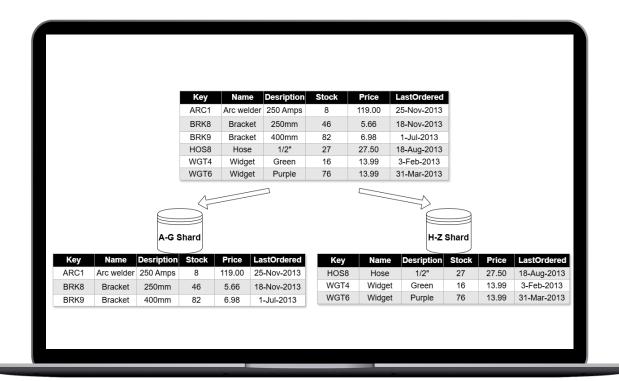


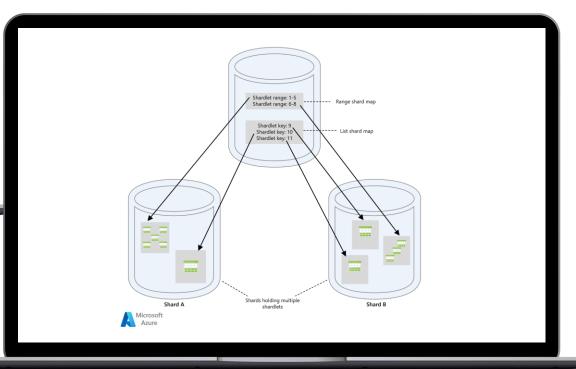
# Partitioning Strategies

Concepts of how to physically divide data into separate data stores

# Horizontal partitioning (sharding)



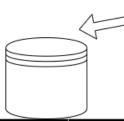




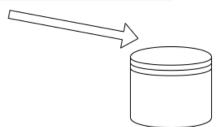




Key	Name	Desription	Stock	Price	LastOrdered
ARC1	Arc welder	250 Amps	8	119.00	25-Nov-2013
BRK8	Bracket	250mm	46	5.66	18-Nov-2013
BRK9	Bracket	400mm	82	6.98	1-Jul-2013
HOS8	Hose	1/2"	27	27.50	18-Aug-2013
WGT4	Widget	Green	16	13.99	3-Feb-2013
WGT6	Widget	Purple	76	13.99	31-Mar-2013



Key	Name	Desription	Price
ARC1	Arc welder	250 Amps	119.00
BRK8	Bracket	250mm	5.66
BRK9	Bracket	400mm	6.98
HOS8	Hose	1/2"	27.50
WGT4	Widget	Green	13.99
WGT6	Widget	Purple	13.99

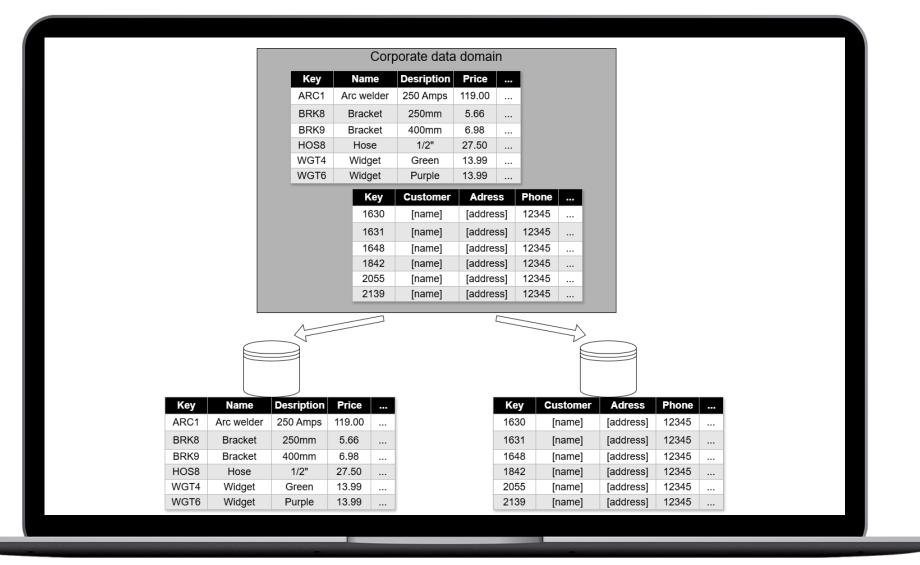


Key	Stock	LastOrdered				
ARC1	8	25-Nov-2013				
BRK8	46	18-Nov-2013				
BRK9	82	1-Jul-2013				
HOS8	27	18-Aug-2013				
WGT4	16	3-Feb-2013				
WGT6	76	31-Mar-2013				

# Functional partitioning







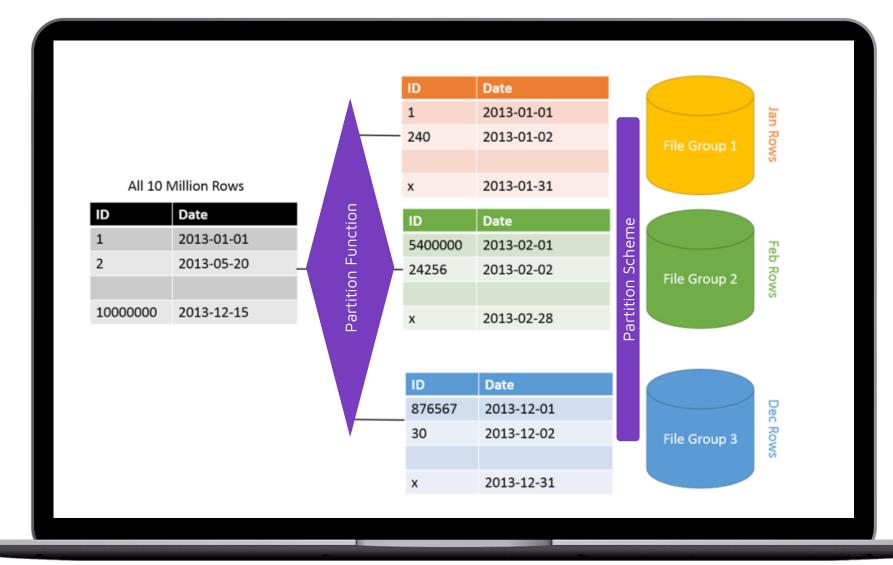


SQL Server deep dive on Partitioning

#### Partitioned tables and indexes







# Demo



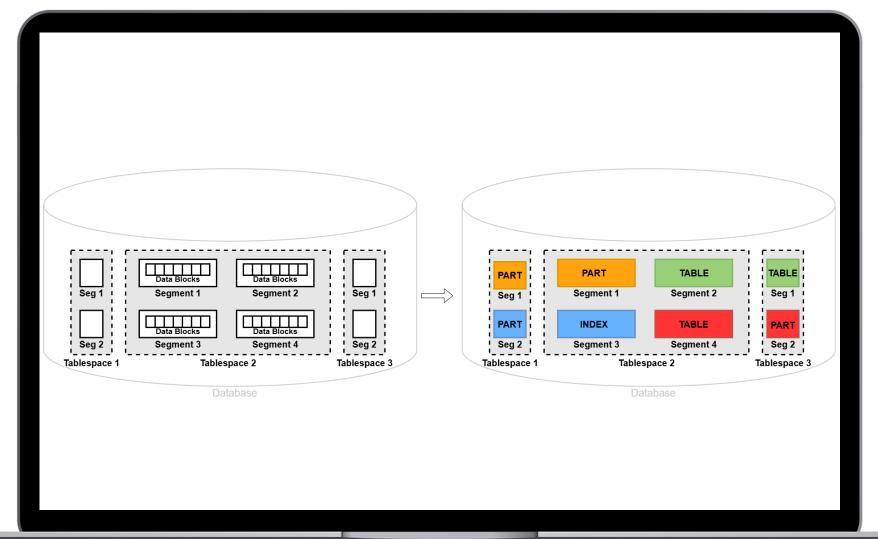
Oracle deep dive on Partitioning





Logical data structures are separate from physical data structures.

Tables, views, indexes, partitions are logical objects. Each object is stored as a separate segment and can have its own storage settings.

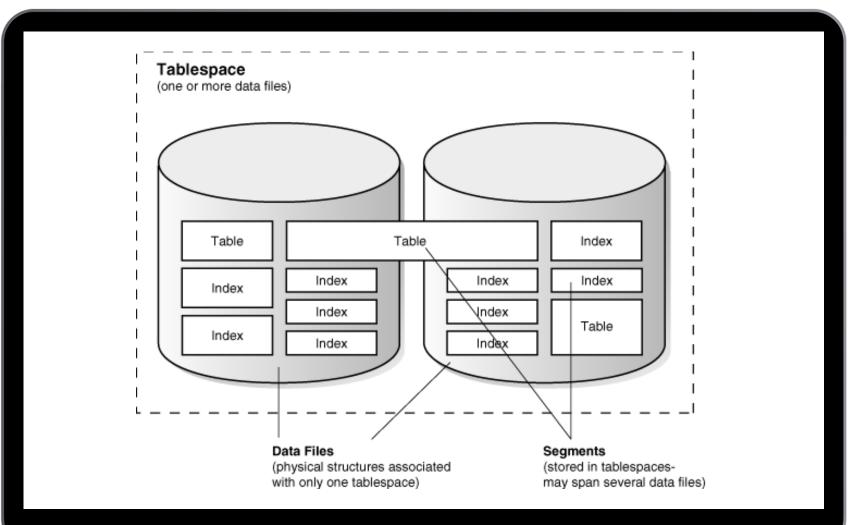






Segments (partitions) may be located in different tablespaces.

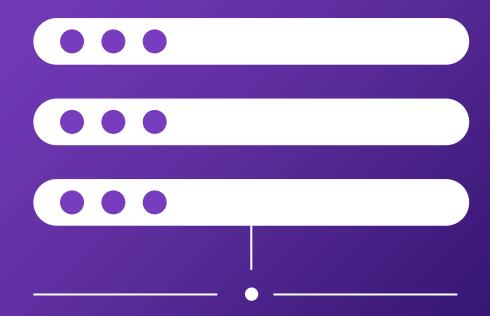
Tablespaces are separate data files or disk groups.
Each partition can reside in a separate storage with its own configuration.



Buffer Cache holds frequently accessed data blocks from any partition. Hot partitions' blocks stay in memory more, effectively prioritizing active data.

Data dictionary stores partition definitions to determine which partition to access. The mapping from a partition key value to the segment is handled by the optimizer.

Parallel scans, assigning different partitions to different query workers, multiple CPU/IO for partitioned table access. Each partition can be scanned independently in parallel, then results combined.

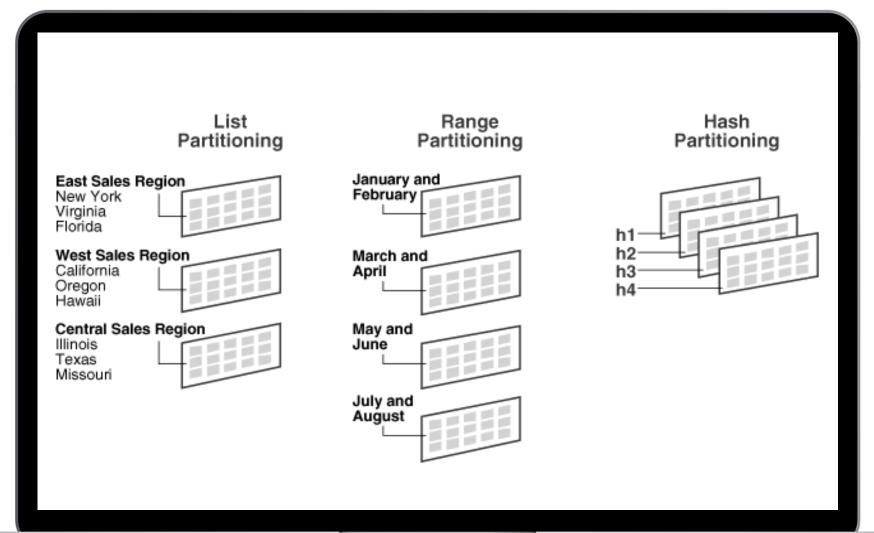


# Partitioning Methods in Oracle

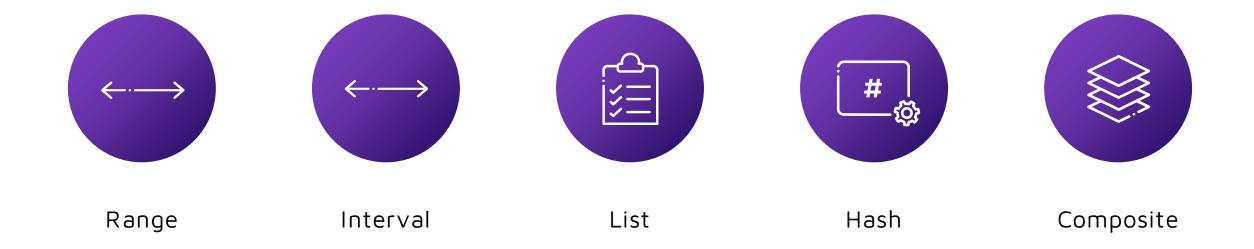




- Range: Data is partitioned by ranges of values. Each partition covers a specific interval of the partition key.
- List: Defined by an explicit list of discrete values for the key.
- Hash: Hash function on the partition key. Rows are evenly distributed across partitions based on hash values.
- Composite: Multi-level partitioning (subpartitions).
   Combining methods.









PostgreSQL deep dive on Partitioning



Partitioned table is a parent virtual table with no data of its own. All data is stored in its child tables - the partitions.

Each partition is an independent table in the storage engine with its own heap file and index files.

ID			Date			
ID	Date	ID	Date		ID	Date
1	2025-01-01	540000	2025-02-01		540000	2025-12-01
240	2025-01-14	24256	2025-02-18		24256	2025-12-10
х	2025-01-31	х	2025-02-28		Х	2025-12-31



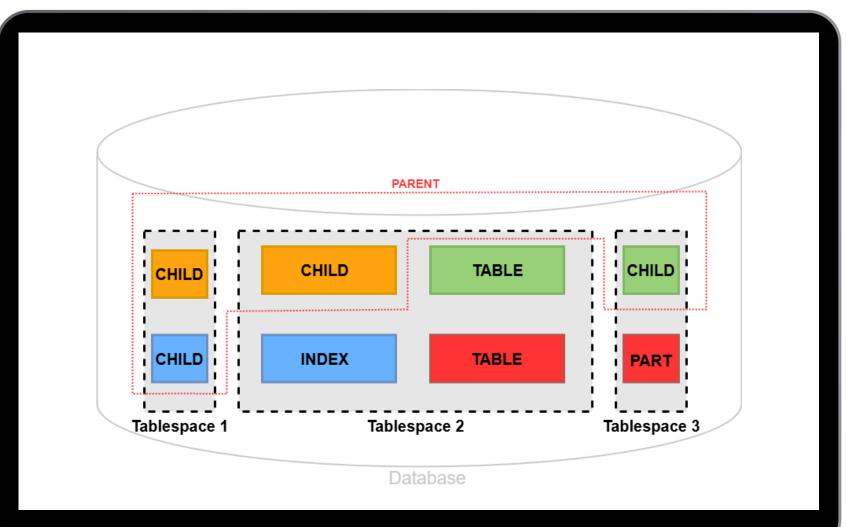


Each partition (child) has its own files on disk.

Parent table stores metadata like partitioning scheme and list of partitions.

Partitions placed in different tablespaces.

Different partitions stored in different file system locations mimicking tiered storage.

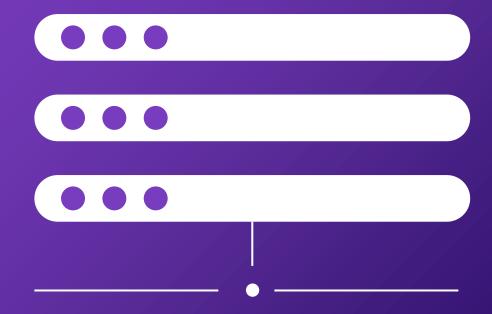


### PostgreSQL: Memory

SQL O Day

Shared buffer pool for caching data. Pages from partition tables are treated the same as any table's pages. Hot partitions' pages will occupy more cache due to frequent access, but there's no dedicated memory per partition.

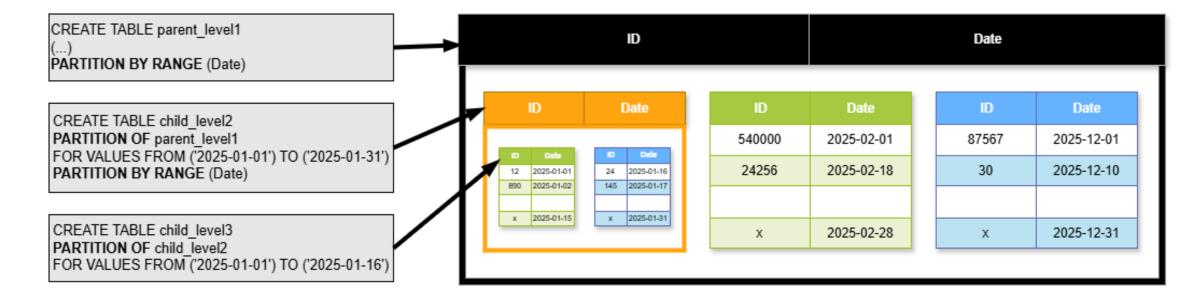
Partition definitions (ranges or lists of values) are stored in the catalog; pg\_class for the child tables, pg\_inherits links them to the parent, pg\_partitioned\_table stores partition key info. Executor routes the inserted row to the appropriate child table. For queries, the planner uses the partition constraints to determine which partitions can be pruned.



## Partitioning in PostgreSQL



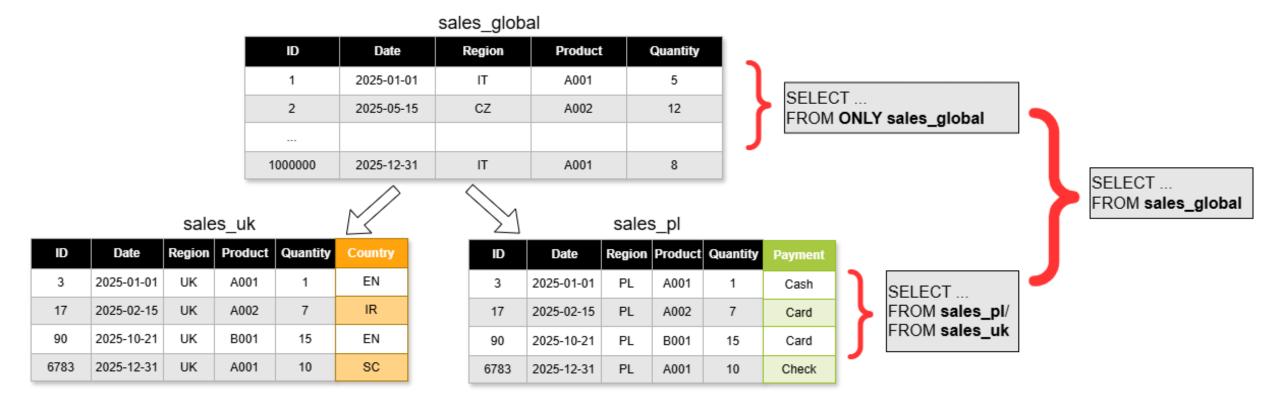
- CREATE TABLE as syntax with one-column partition specified. Individual names.
- Partitioning existing table is not possible.
   Attaching/detaching existing table as a partition is possible.
- Multi-level partitioning. 3 levels for best performance.
- No indexes possible at TABLE\_PARENT, only at TABLE\_CHILD (leaf).
- INSERT/SELECT/DELETE possible on TABLE\_PARENT and TABLE\_CHILD



# Partitioning in PostgreSQL



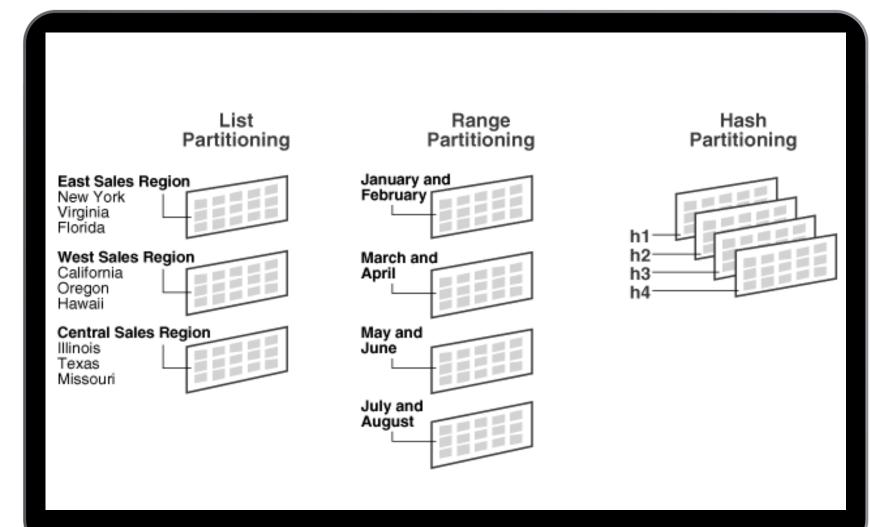
Inheritance: Legacy Method where partitioning is done via table inheritance and check constraints manually. This involved parent table + child tables and triggers or application logic for routing inserts. Declarative partitioning automates this.







- Range: Data is partitioned by ranges of values. Each partition covers a specific interval of the partition key.
- List: Defined by an explicit list of discrete values for the key.
- Hash: Hash function on the partition key. Rows are evenly distributed across partitions based on hash values.
- Multi-level: Nested child tables.
   Each table can have different method.









Range



List





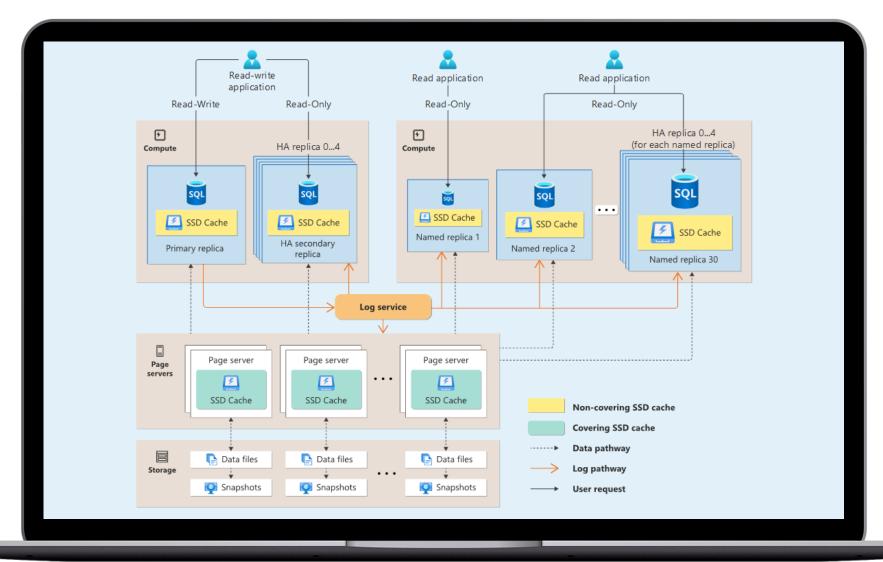


Multi-level



# What you need to know about partitioning Azure SQL Database

On Azure SQL DB Hyperscale example





#### Table partitioning and table maintenance

• Easier index and statistics maintenance on large tables

#### Table partitioning and columnstore indexes

• Each partition has to contain a minimum of 1 million rows for optimal compression and performance of clustered columnstore tables

#### Table partitioning and compression

- Implement a table partitioning strategy first, and then implement your compression strategy
- Possibility of using compressions for different sections of the partitioned table

#### Table partitioning and tempdb

- When implementing table partitioning, some data may have to be temporarily stored in tempdb. tempdb space is limited, so filling the tempdb completely, will be resulting in a failure.
- Implement the data copy from source table to destination partitioned table in batches



#### Table partitioning and Index alignment

- Index is implemented on top of the same partition scheme as its base table.
- Partition switch operations qualify as being metadata-only operations

#### Table partitioning and Replication

- There are requirements and limitations in terms of partition functions and partitions in terms of replication
- Manual administration effort might be needed at subscribers

#### Table partitioning and statistics

 Possibility to use incremental statistics: computed at partition level

#### Table partitioning and Change data capture (CDC)

- Doing operations in a partitioned table can generate inconsistencies if the partitioned table is enabled for change data capture
- Those operations are PARTITION SWITCH, MERGE and SPLIT.



#### Partitioning in place

- The movement of data among partitions happens in-place inside the same table, making this method less space consuming as there is no duplication of tables
- Data movement inside the existing table limits the options to stop and roll back the partitioning operation if desired.
- During the data movement, locks are held on the pages of the table that could block other activities, making it difficult to do online.

#### Partitioning heap tables

• For a heap, a clustered index over a partition scheme has to be created to actually shuffle the rows among the partitions, and subsequently dropped so the table becomes a heap table again.

#### Partitioning into a new table

- Original table is available during process
- Possibility to start/stop/rollback if needed
- Can be used for testing to find best partitioning configuration
- Higher complexity doing the data synchronization in comparison with in-place partitioning

#### Sliding window table partitioning scenario

 Data management scenario for a table in which the same number of partitions is kept over time, and as new time periods start, new partitions are created and brought into the table while the older partitions are evicted



		Oracle	SQL Server	PostgreSQL	
		Yes	Yes		
		defining upper boundary	boundaries for inner	Yes	
		only, right edge partition	partitions defined, edge	edge partitions open	
Types	Range	open (ptional)	partition open (required)	(optional)	
	Interval	Yes	No	No	
	List	List with subpartitions	No	List without subpartitions	
	Hash	Yes	No	Yes	
	Composite	Yes	No	Yes	
Subpartitioning	3	Yes (2 levels)	No (1 level)	Yes (multilevel)	
				Yes, inheritance. On individual	
Indexes		Yes, fully suported	Yes, fully suported	partitions, not the parent	
Partitioning existing tables		No	No	No	
Partition prunning		built-in	built-in	setup enabled/disabled	
Separate storage		Yes (on prem)	Yes (on prem)	Yes (on prem)	
Parallelism		Yes	Yes	Yes	

# lingaro