

# 17th edition SQLDay Conference

12-14 May 2025, WROCŁAW + ONLINE



---

## Platinum sponsors

---



DATUMO



Woodler

---

## Gold sponsors

---

lingaro



TECHNOLOGY  
INNOVATION  
DATA  
KNOWLEDGE

hidk

software  
**one**

 Dataedo

---

## Silver sponsors

---

Capgemini



elitmind

&F

 kursy  
FABRIC

# Going Live with dbt-core on Databricks and MS Fabric

SQLDay Wrocław 2025-05-14

# Speaker



## Tomasz Kostyrka, PL

- Data Platform Architect @GetInData | Part of Xebia
- ~12 years in Data
- Azure/Databricks/PowerPoint
- Databricks Architect Champion
  
- <https://www.linkedin.com/in/tomasz-kostyrka/>
- <https://sessionize.com/tomasz-kostyrka/>
- <https://pl.seequality.net/>



# Session



## Plan:

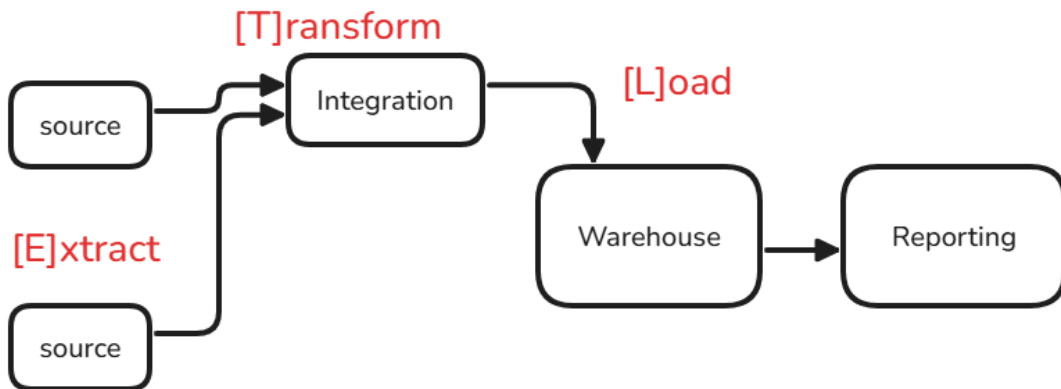
- extremely quick introduction to dbt
- Go Prod!
- Databricks/Microsoft Fabric
- Documentation

## Disclaimers:

- Opinions are personal
- I'll oversimplify a lot
- Let's have some fun

# **extremely quick introduction to dbt**

# Extract-Transform-Load



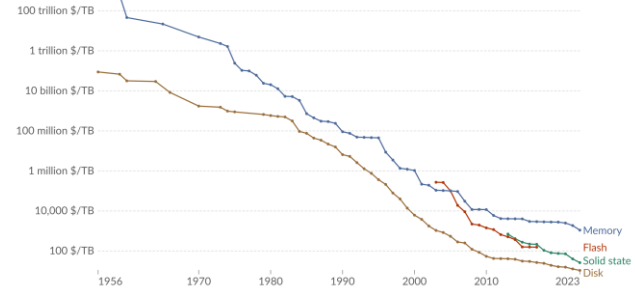
- SQL Server Integration Services (DataStage, AbInitio, Talend)
- SQL Server (Oracle, Teradata, Exadata)
- SQL Server Reporting Services (Qlik, Tableau)



## Historical price of computer memory and storage

This data is expressed in US dollars per terabyte (TB), adjusted for inflation. "Memory" refers to random access memory (RAM), "disk" to magnetic storage, "flash" to special memory used for rapid data access and rewriting, and "solid state" to solid-state drives (SSDs).

Our World  
In Data



Data source: John C. McCallum (2023); U.S. Bureau of Labor Statistics (2024)

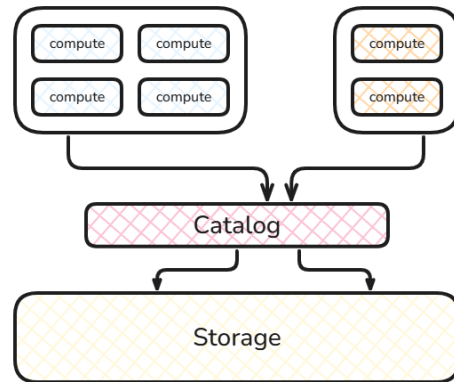
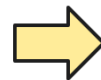
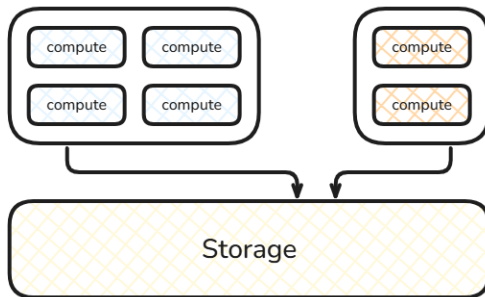
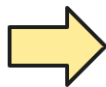
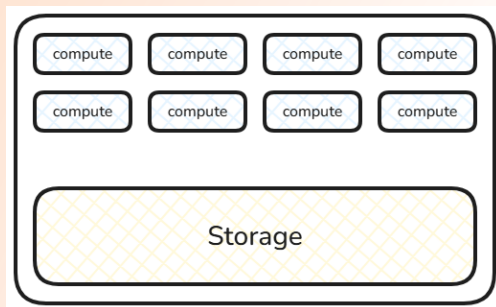
OurWorldInData.org/technological-change | CC BY

Note: For each year, the time series shows the cheapest historical price recorded until that year. This data is expressed in constant 2020 US\$.



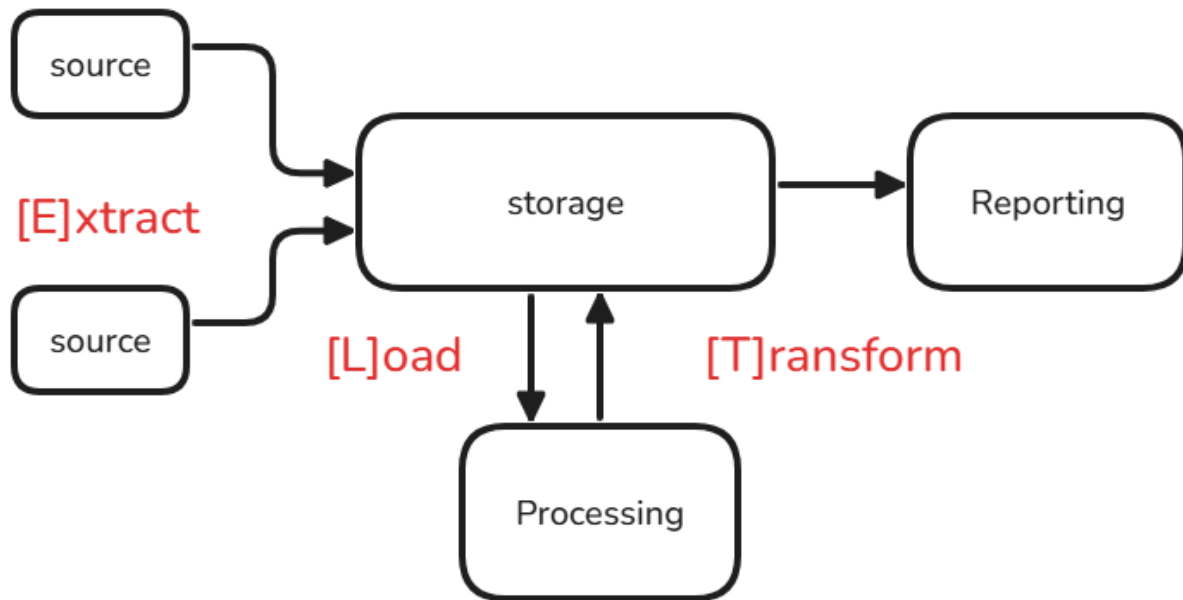
Google Cloud

# From Warehouses to Lakehouses in 30 seconds

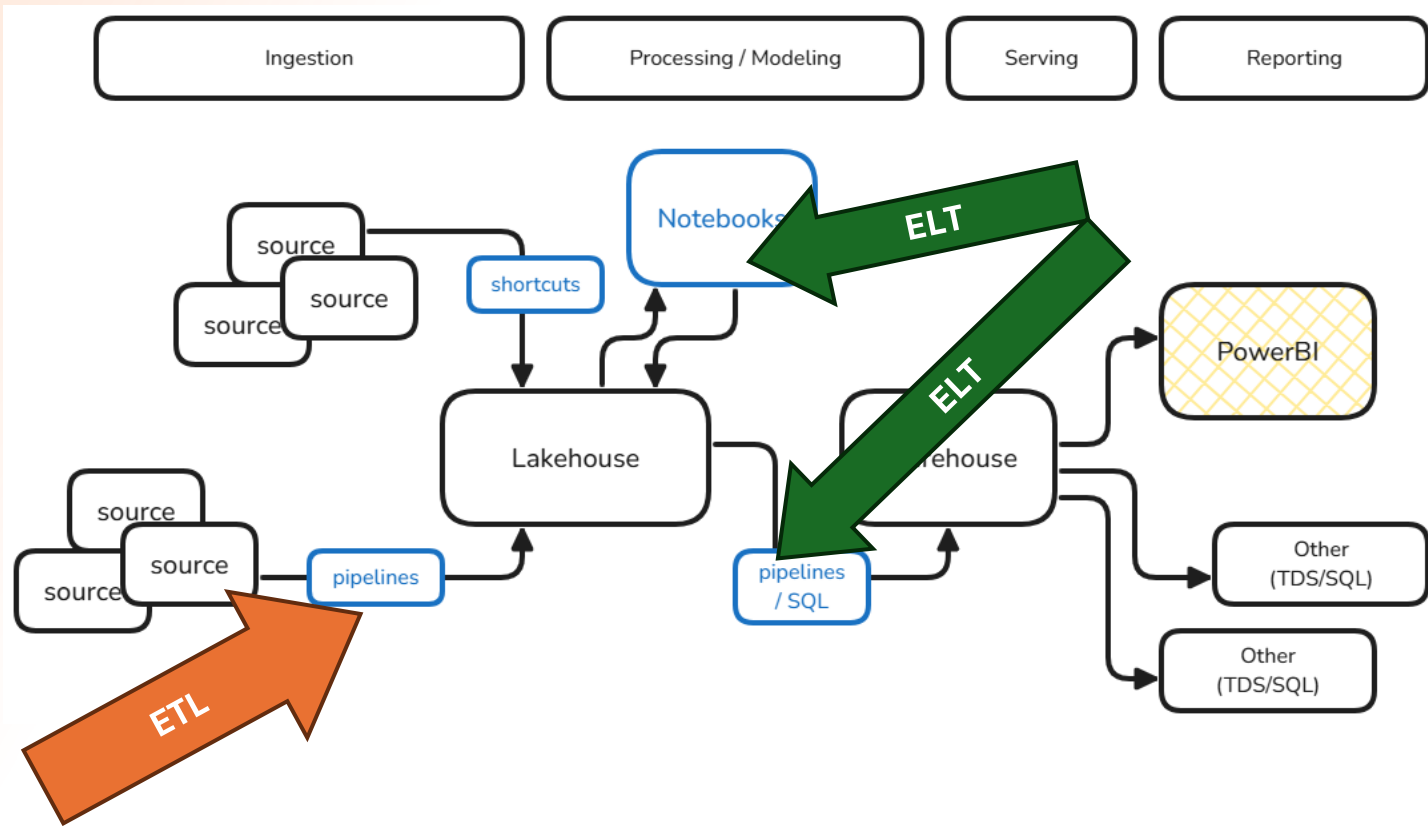




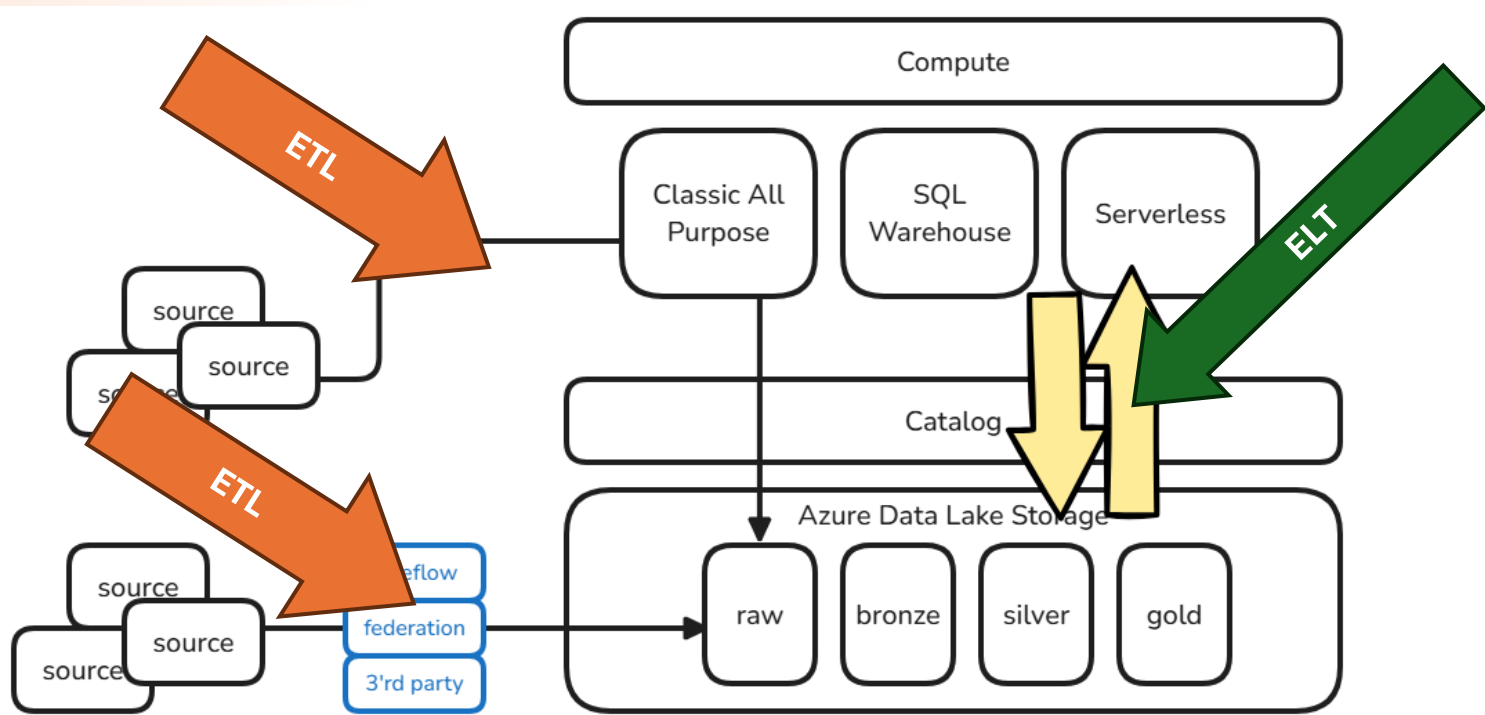
# Extract-Load-Transform



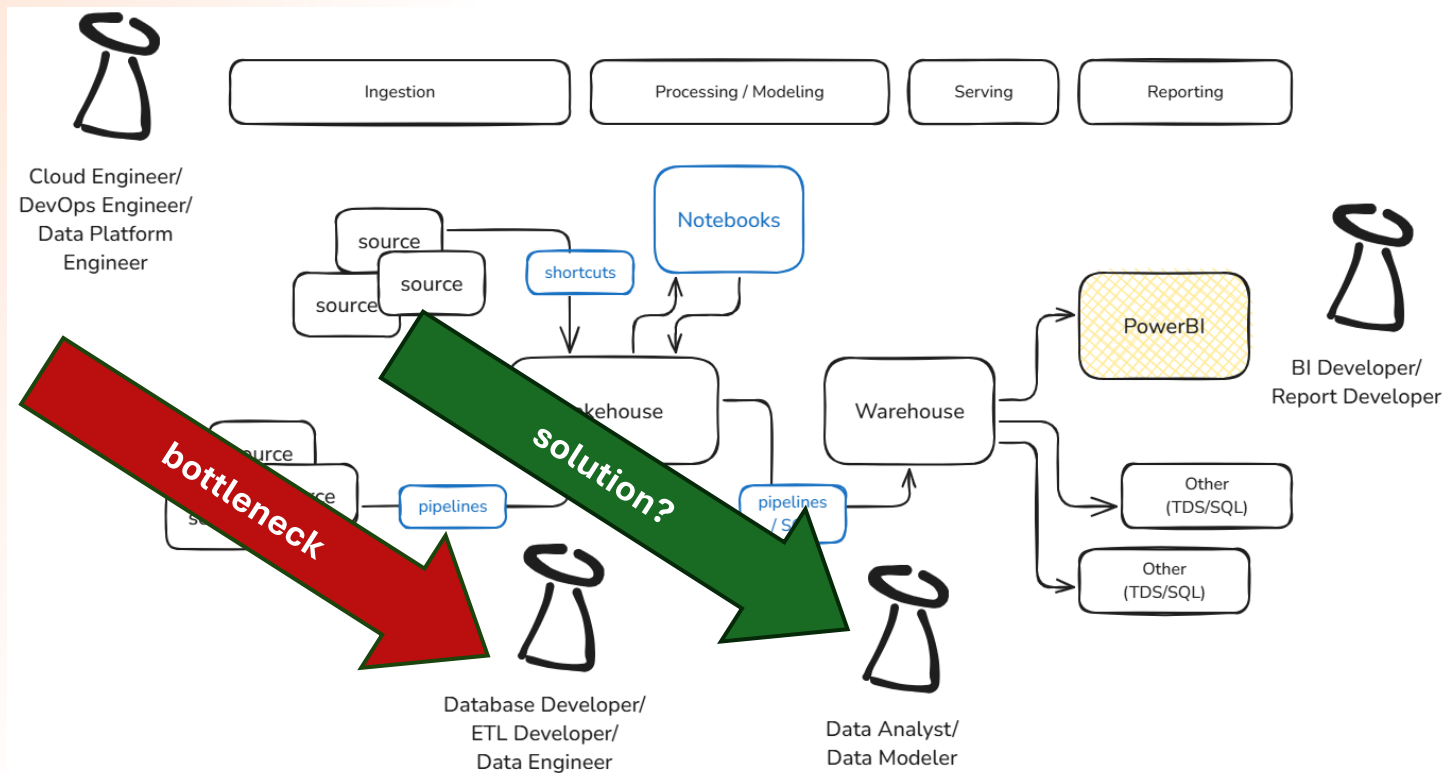
# Microsoft Fabric

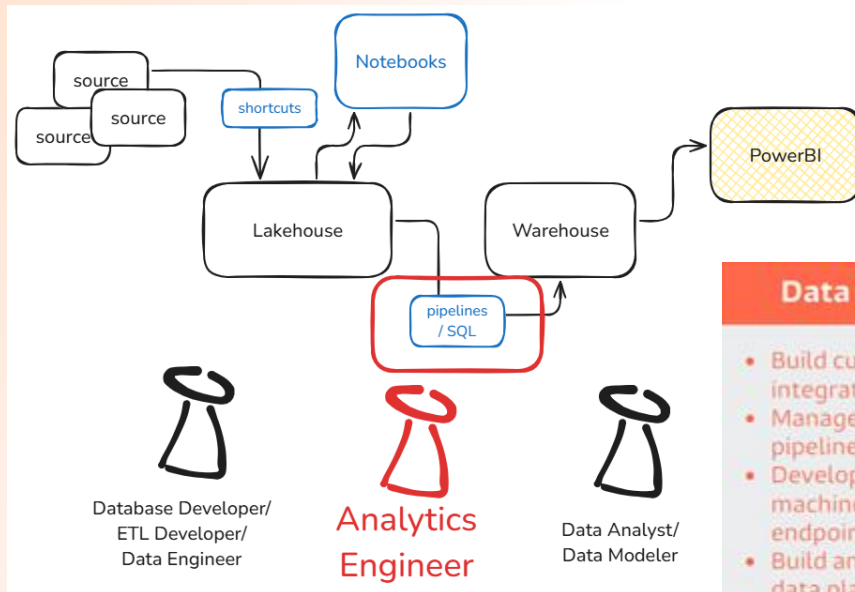


# Databricks



# Roles





### Data Engineer

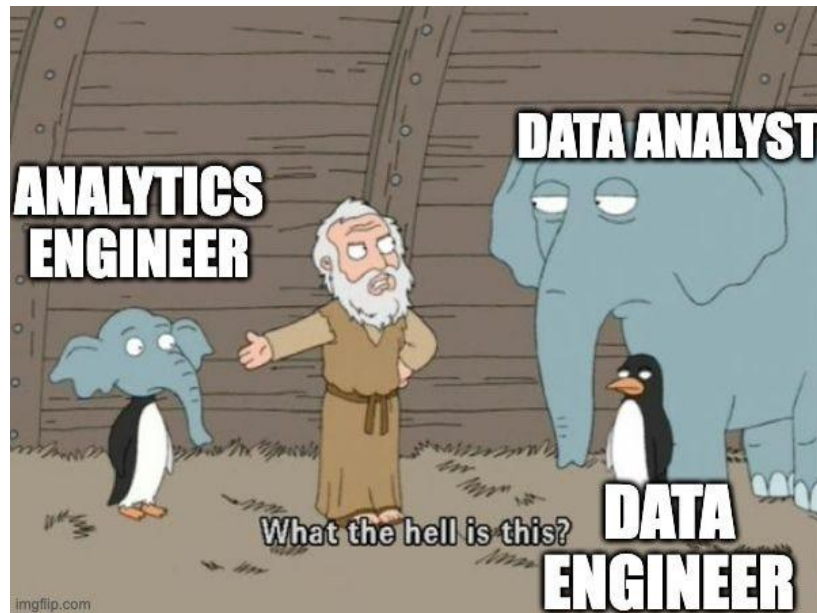
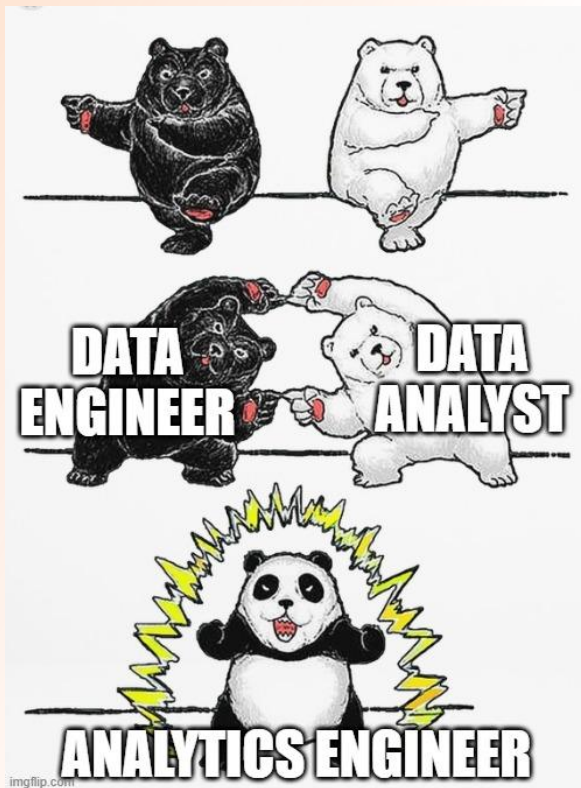
- Build custom data integrations
- Manage overall pipeline orchestration
- Develop & deploy machine learning endpoints
- Build and maintain the data platform
- Data warehouse performance optimizations

### Analytics Engineer

- Provide clean, transformed data ready for analysis
- Apply software engineering best practices to analytics code (ex: version control, testing, continuous integration)
- Maintain data documentation & definitions
- Train business users on how to use data visualization tools

### Data Analyst

- Deep insights work (ex: why did churn spike last month? what are the best acquisition channels?)
- Work with business users to understand data requirements
- Build critical dashboards
- Forecasting





dbt™

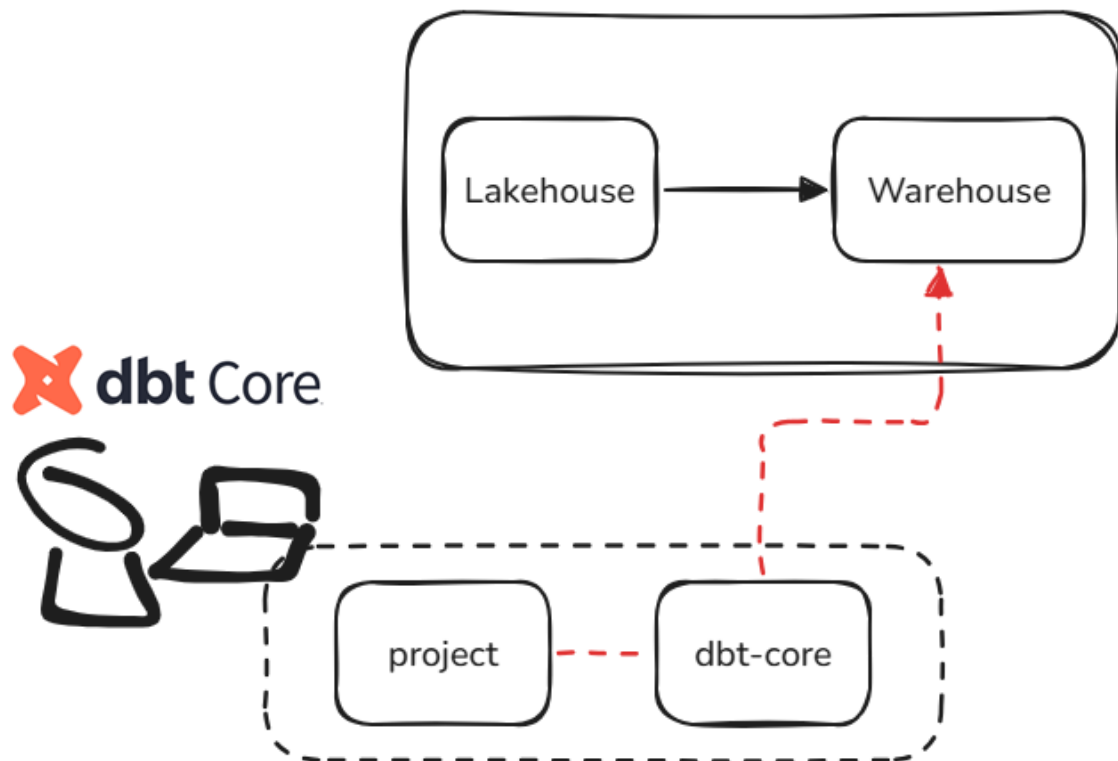


coalesce®



SQLMesh

# Local





# Go Prod!



**dbt Cloud**



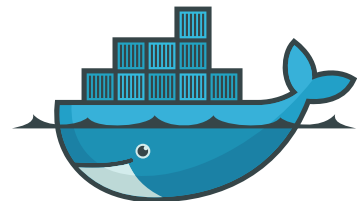
**dbt Core**



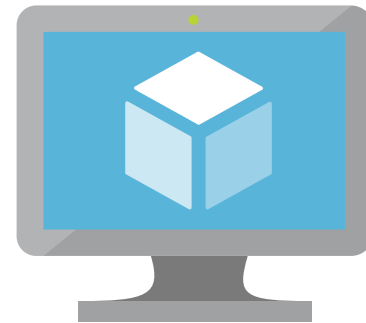
# Host

We need a **compute environment**.

- A developer's laptop.
- A virtual machine (e.g., on-prem or cloud VM).
- A container (e.g., Docker).



docker



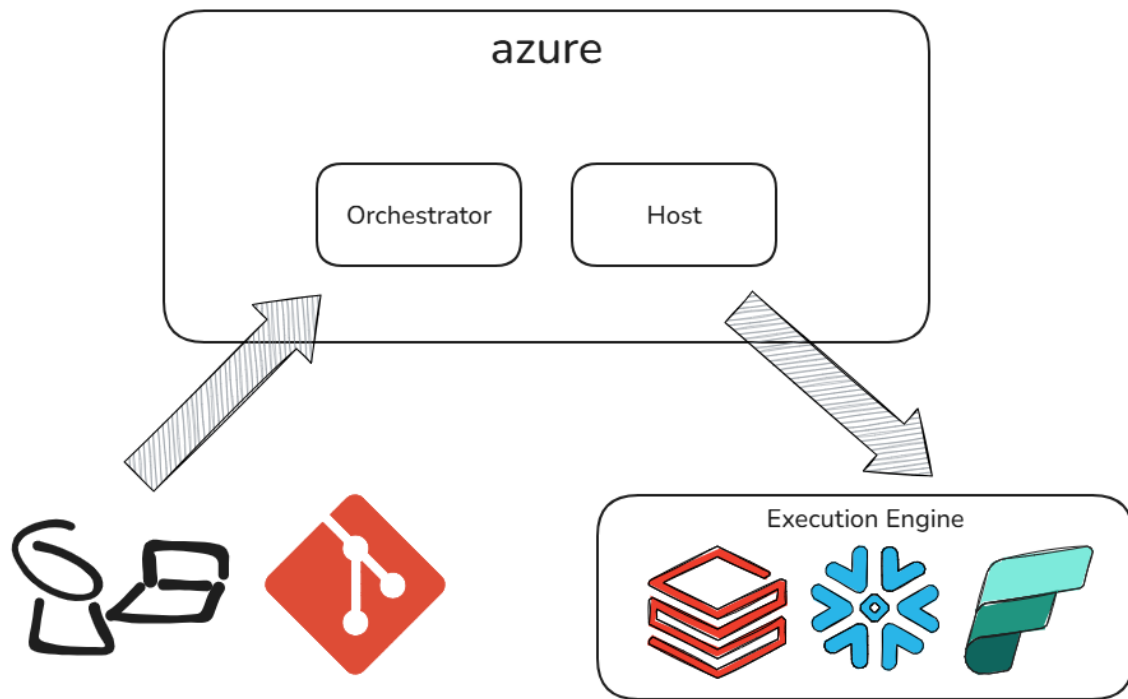
# Scheduler

We need a **scheduler**.

- cron jobs
- Apache Airflow
- Azure Data Factory
- Databricks Workflows

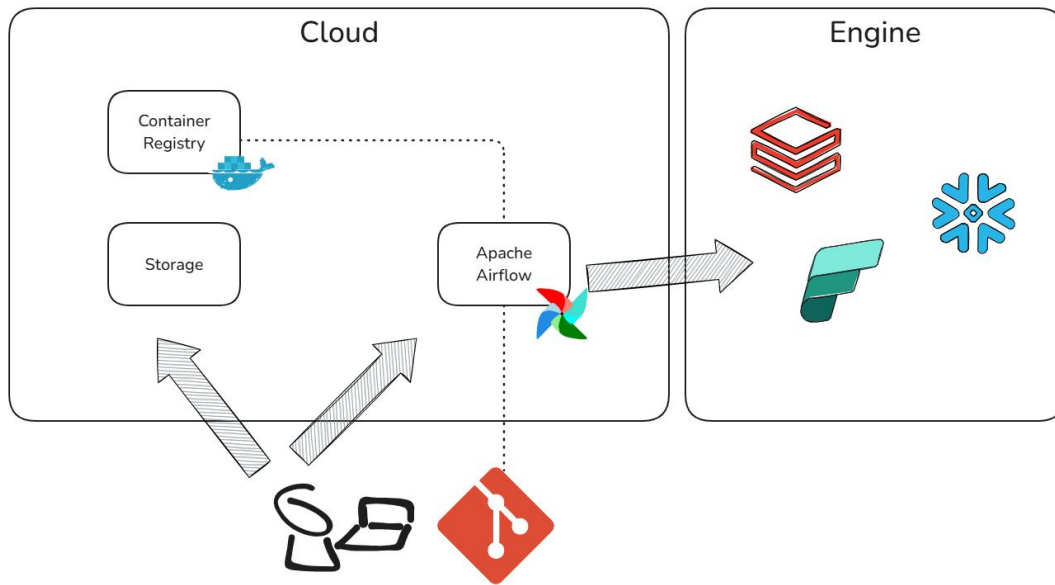


Apache  
**Airflow**



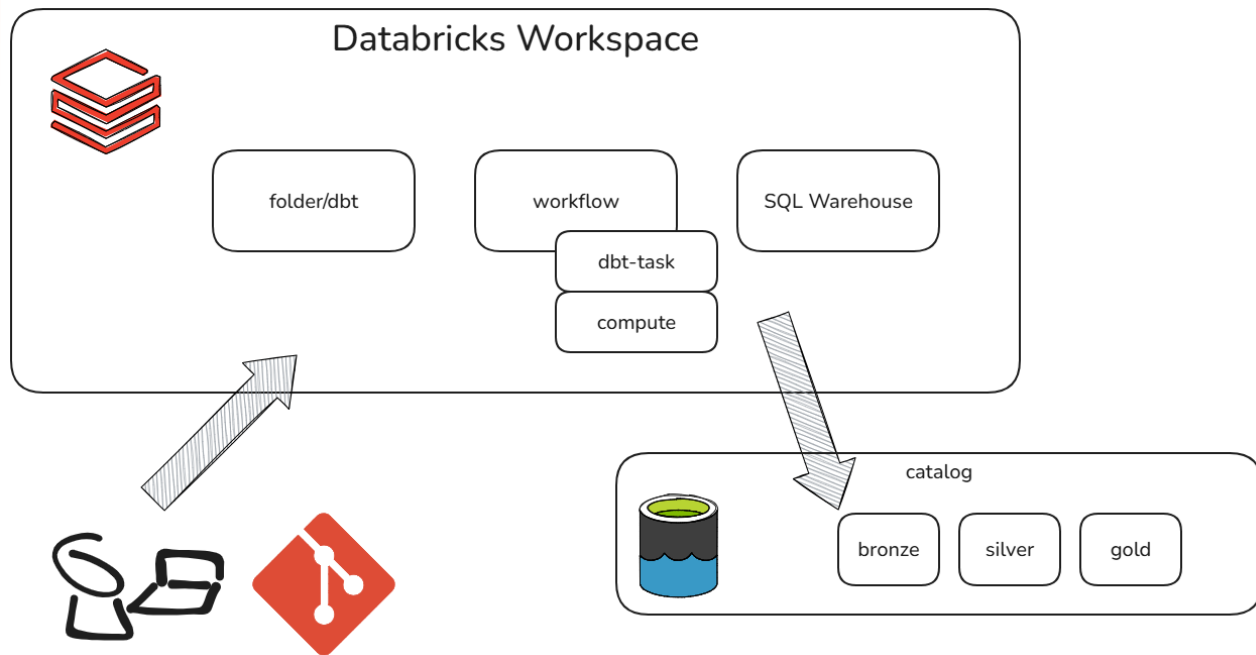
# Apache Airflow

- Astronomer
- Google Cloud Composer
- Amazon MWAA (Managed Workflows for Apache Airflow)
- Azure Data Factory with Airflow integration



# Databricks





Runs Tasks

dbt-run

3 dbt commands

Task name\*

dbt-run

Type\*

dbt

Source\*

Workspace

Project directory\*

/Workspace/Users/tomasz.kostyrka@getindata.com

dbt commands\*

1

dbt deps

{ }

2

dbt seed

{ }

✕

3

dbt run

{ }

✕

[Add](#)

SQL warehouse\*

2xsmall-tkostyrka (2XS)

Warehouse catalog\*

sqlday2025

Warehouse schema\*

dbo

dbt CLI compute\*

Serverless

Autoscaling

Environment and Libraries\*

Default

Cancel

Create task



## Job details

Job ID

159261289714592

Creator

Tomek Kostyrka (GID)

Run as

Tomek Kostyrka (GID)

Tags

[Add tag](#)

Description

[Add description](#)

Lineage

No lineage information for this job.  
[Learn more](#)

## Schedules &amp; Triggers

None

[Add trigger](#)

## Job parameters

No job parameters are defined for this job

[Edit parameters](#)

## Job notifications

No notifications

[Edit notifications](#)

Duration and streaming backlog thresholds

No thresholds defined

[Add metric thresholds](#)

## Permissions



EXPLORER

## ▼ DATABRICKS

## ▼ bundle

&gt; targets

## ▼ workflows

! dbxdemo\_dbt.yml

! variables.yml

## ▼ dbt

&gt; analyses

&gt; logs

&gt; macros

&gt; models

&gt; seeds

&gt; snapshots

&gt; target

&gt; tests

! .user.yml

! dbt\_project.yml

! profiles.yml

&gt; notebooks

! databricks.yml

! dbxdemo\_dbt.yml X

bundle &gt; workflows &gt; ! dbxdemo\_dbt.yml &gt; {} resources &gt; {} jobs &gt; {} dbxdemo\_dbt &gt; [ ] tasks &gt; {} 0 &gt; [ ] libraries &gt; {} 0 &gt;

1 resources:

2 jobs:

3 dbxdemo\_dbt:

4 name: dbxdemo\_dbt

5

6 queue:

7 enabled: true

8

9 schedule:

10 quartz\_cron\_expression: 0 30 0 \* \* ?

11 timezone\_id: UTC

12 pause\_status: PAUSED

13

14 email\_notifications:

15 on\_failure:

16 - tomasz.kostyrka@getindata.com

17

18 tasks:

19 - task\_key: gold

20 dbt\_task:

21 project\_directory: \${workspace.root\_path}/files/dbt

22 profiles\_directory: .

23 commands:

24 - 'dbt deps -t \${bundle.target}'

25 - 'dbt seed -t \${bundle.target}'

26 - 'dbt run -t \${bundle.target}'

27 libraries:

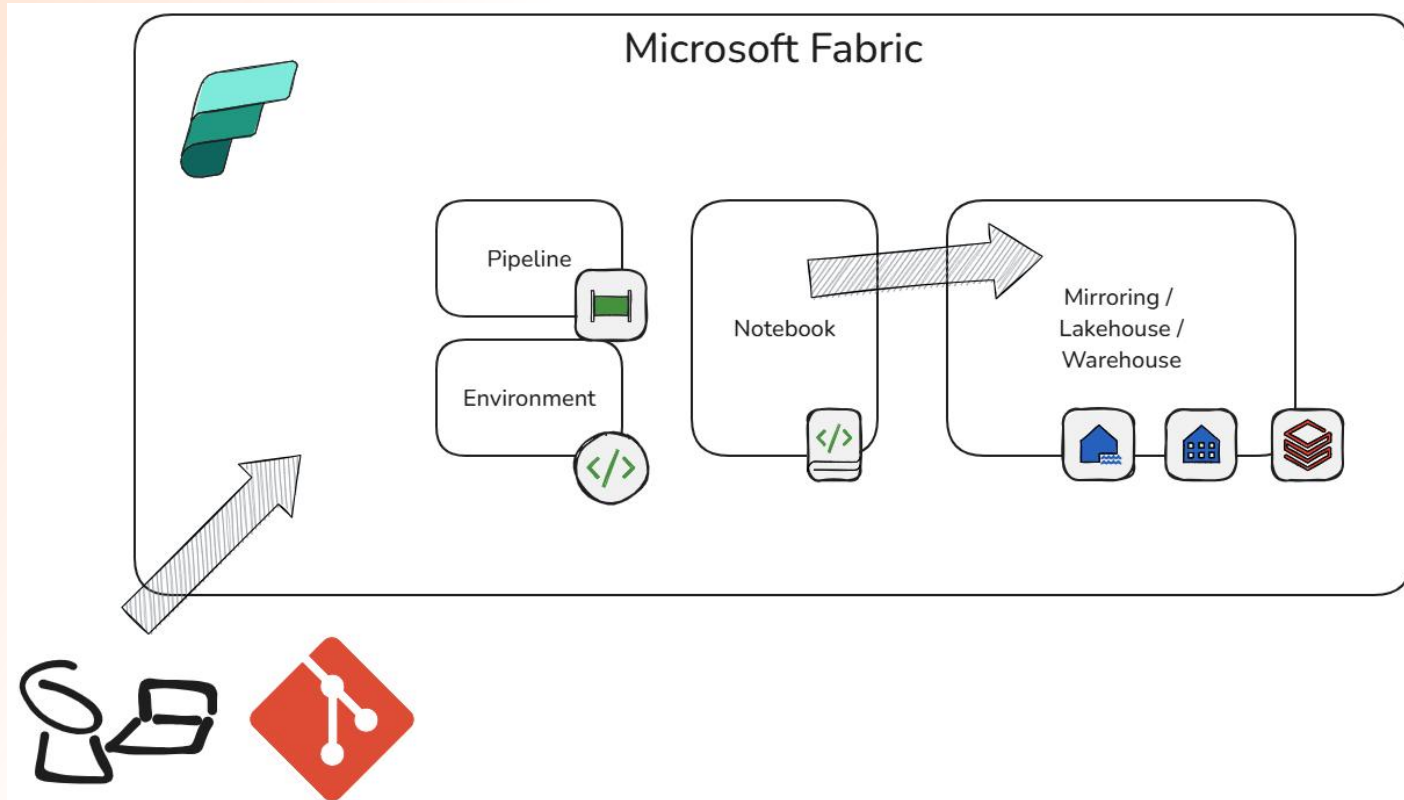
28 - pypi:

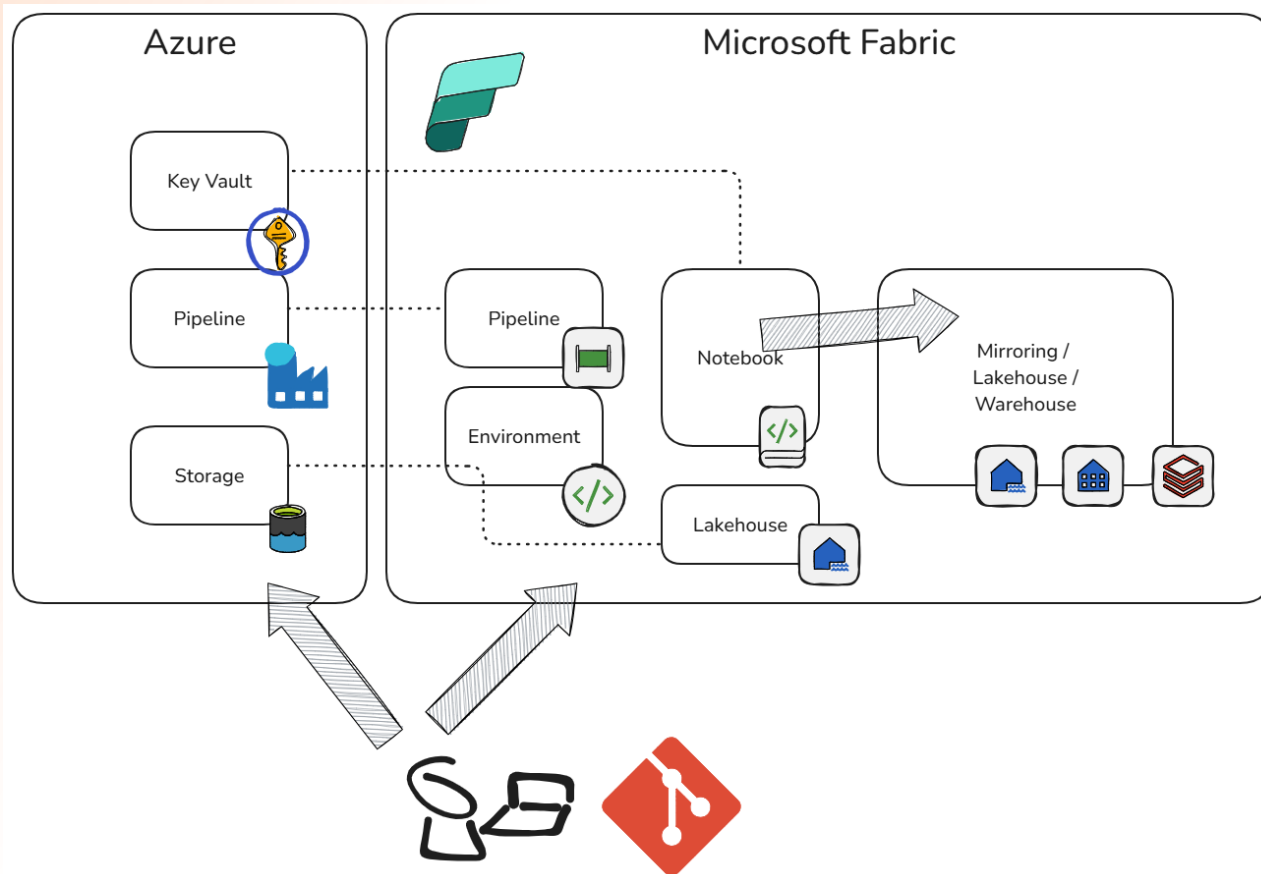
29 package: dbt-databricks&gt;=1.8.0,&lt;2.0.0

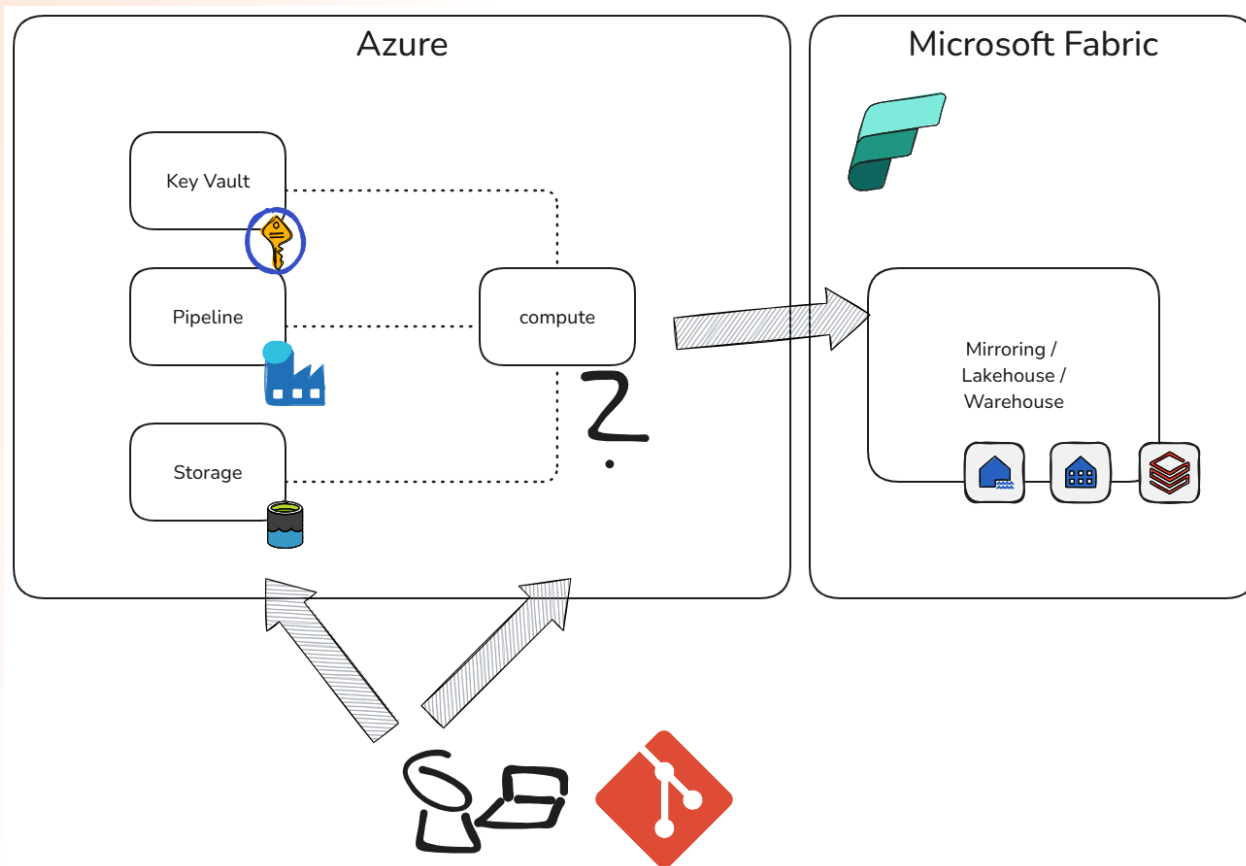


```
33 strategy:
34   runOnce:
35     deploy:
36       steps:
37         - checkout: self
38
39         - task: UsePythonVersion@0
40           inputs:
41             versionSpec: '3.11'
42             displayName: 'Use Python 3.11'
43
44         - script: |
45             curl -fsSL https://raw.githubusercontent.com/databricks/setup-cli/main/install.sh | sh
46             databricks -v
47             displayName: install databricks
48
49         - script: |
50             curl -sSL https://install.python-poetry.org | python3 -
51             poetry --version
52             displayName: install poetry
53
54         - task: AzureCLI@2
55           displayName: deploy bundle
56           inputs:
57             azureSubscription: '${{ parameters.serviceconn }}'
58             scriptType: bash
59             scriptLocation: inlineScript
60             addSpnToEnvironment: true
61             workingDirectory: '${{ parameters.workingDirectory }}'
62             inlineScript: |
63               export ARM_CLIENT_ID=$servicePrincipalId
64               export ARM_CLIENT_SECRET=$servicePrincipalKey
65               export ARM_TENANT_ID=$tenantId
66
67               databricks bundle deploy -t '${{ parameters.target }}
```

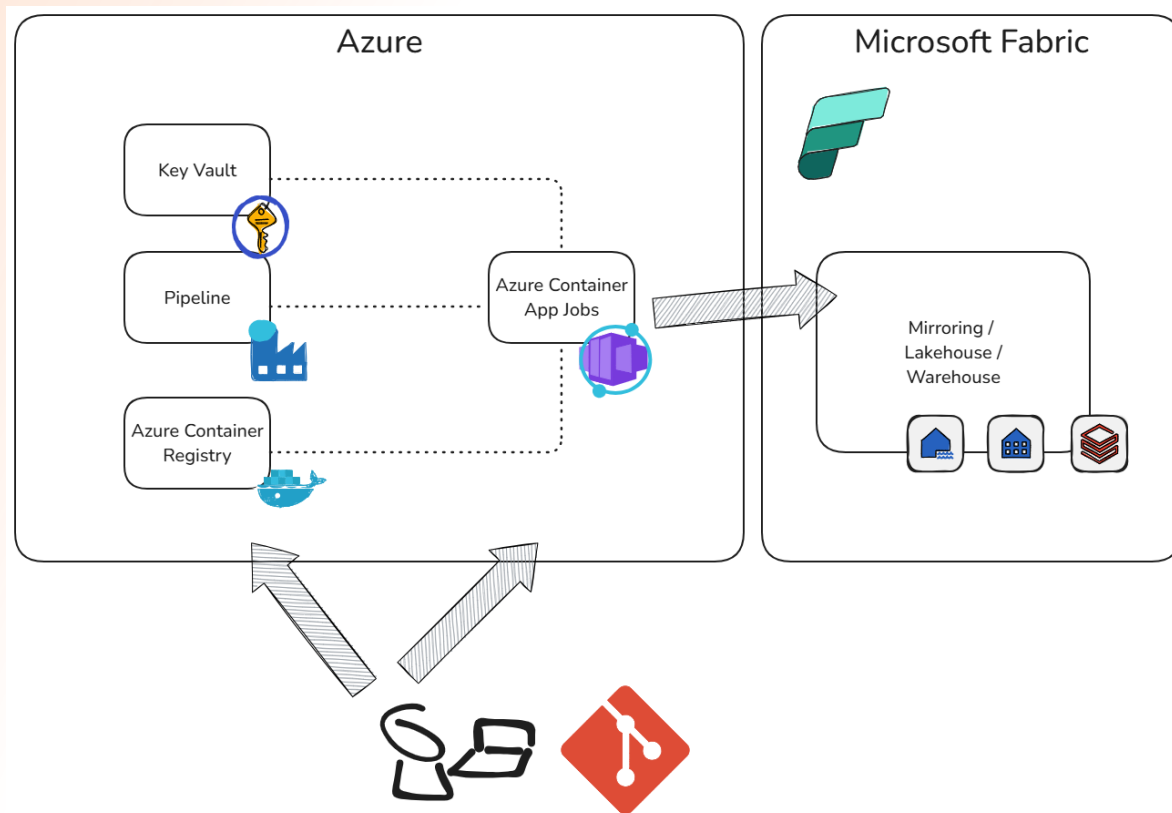
# Microsoft Fabric









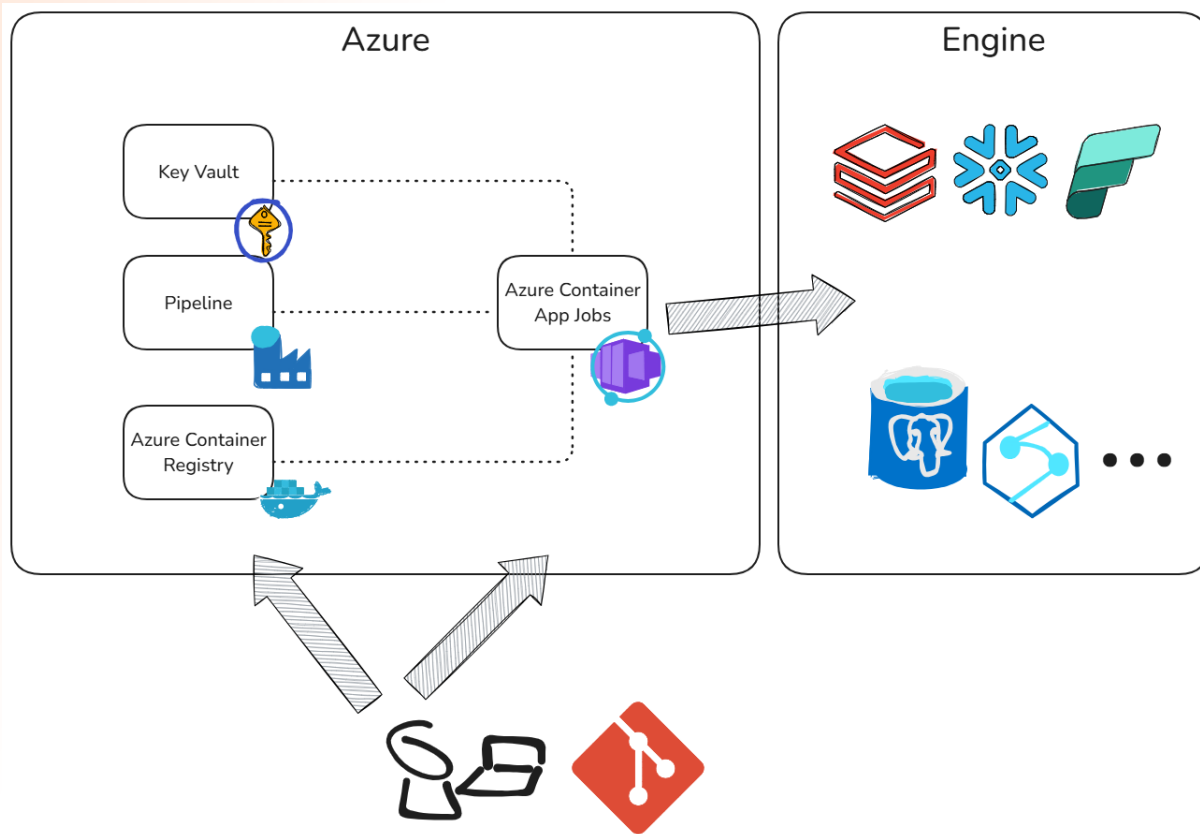


# Dockerfile



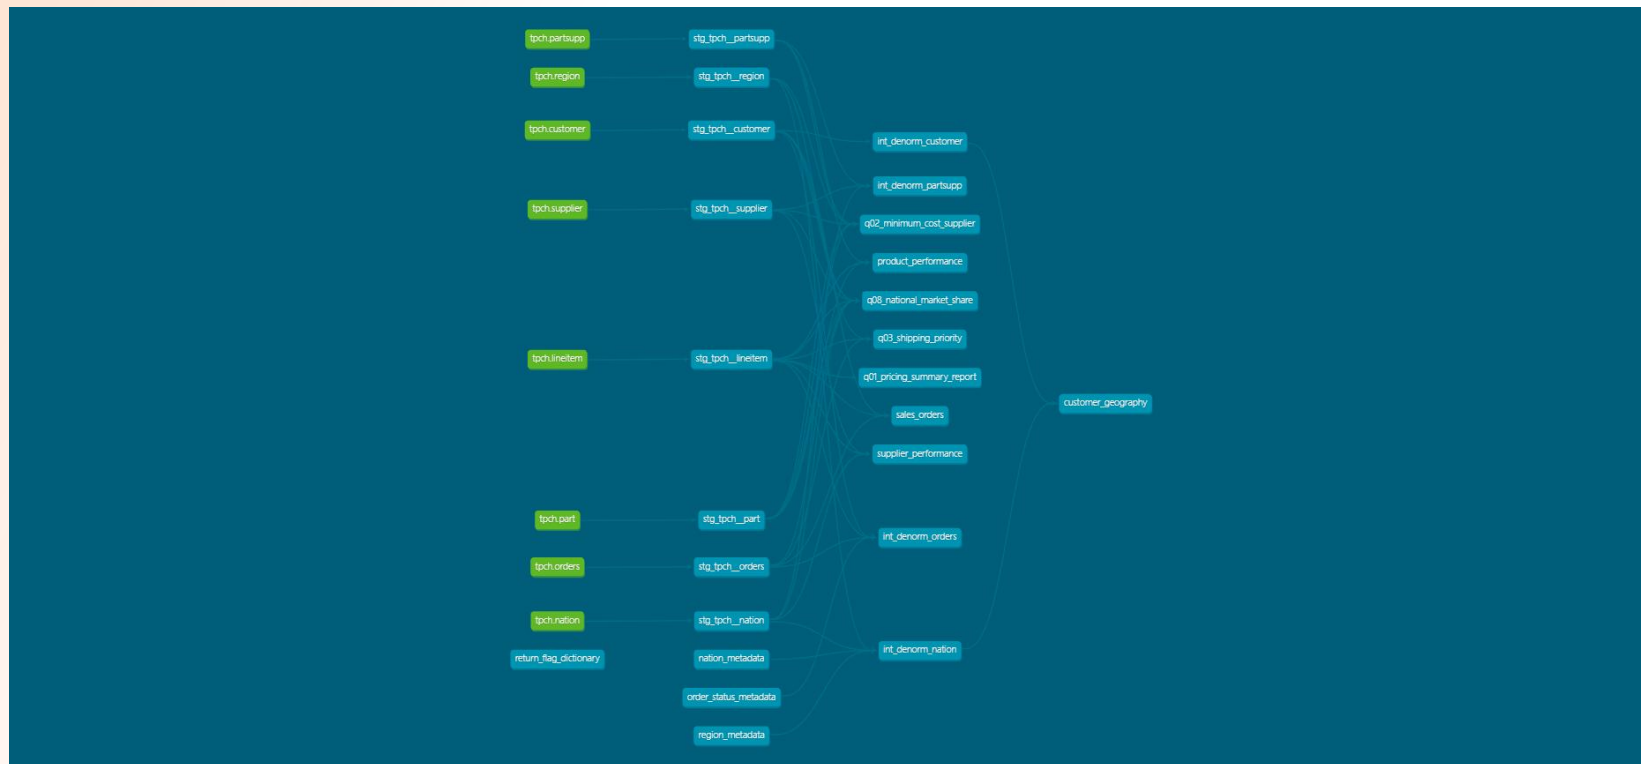
```
1 FROM ghcr.io/dbt-labs/dbt-core:1.9.2
2
3 RUN apt-get update && \
4     apt-get install -y --no-install-recommends unixodbc-dev curl gnupg2 && \
5     rm -rf /var/lib/apt/lists/*
6
7 RUN curl https://packages.microsoft.com/keys/microsoft.asc | tee /etc/apt/trusted.gpg.d/microsoft.asc && \
8     curl https://packages.microsoft.com/config/debian/11/prod.list | tee /etc/apt/sources.list.d/mssql-release.list && \
9     apt-get update && \
10    ACCEPT_EULA=Y apt-get install -y msodbcsql18 && \
11    apt-get clean && \
12    rm -rf /var/lib/apt/lists/*
13
14 RUN pip install dbt-fabric~=1.9.2
15
16 COPY . /usr/app
17 COPY profiles.yml /root/.dbt/
18
19 RUN dbt deps
20 CMD ["run"]
```

```
61 # Download Artifact
62 # -----
63 - task: DownloadPipelineArtifact@2
64   inputs:
65     buildType: current
66     artifactName: ${ parameters.artifactname }
67     targetPath: $(Pipeline.Workspace)/drop/dbt
68
69 # Download Artifact
70 # -----
71 - task: Docker@2
72   displayName: buildAndPush to ACR
73   inputs:
74     command: buildAndPush
75     repository: ${ parameters.repository }
76     dockerfile: $(Pipeline.Workspace)/drop/dbt/Dockerfile
77     containerRegistry: ${ parameters.registry }
78     tags: |
79       latest
80       $(Build.BuildId)
```



# Documentation

# Lineage



# Documentation



The screenshot shows the dbt Docs web interface in a browser. The URL is `localhost:8080/#/model/model.project.q01_pricing_summary_report#columns`. The interface includes a sidebar on the left with a tree view of the project structure, a top navigation bar with tabs for 'Project', 'Database', and 'Group', and a main content area on the right.

**dbt**

Search for models...

**Overview**

Project Database Group

**Sources**

- tpch

**Projects**

- project
  - models
    - intermediate
      - int\_denorm\_customer
      - int\_denorm\_nation
      - int\_denorm\_orders
      - int\_denorm\_partsupp
    - marts
      - customer\_geography
      - product\_performance
      - sales\_orders
      - supplier\_performance
    - reports
      - q01\_pricing\_summary\_report**
      - q02\_minimum\_cost\_supplier
      - q03\_shipping\_priority
      - q08\_national\_market\_share
  - staging
  - seeds
  - tests

**q01\_pricing\_summary\_report** view

Details Description Columns Depends On Code

**Columns**

COLUMN	TYPE	DESCRIPTION	CONSTRAINTS	DATA TESTS	MORE?
return_flag		Indicates whether a returned item is part of the ord...			>
line_status		Status of the order line (e.g., 'O' for open, 'F' for fill...			>
sum_qty		Total quantity of items ordered for the given return...			>
sum_base_price		Sum of the extended price (base sales amount) bef...			>
sum_disc_price		Sum of the price after discount (i.e., extended_price...			>
sum_charge		Sum of the total charge (discounted price * (1 + ta...			>
avg_qty		Average quantity per line item.			>
avg_price		Average extended price per line item.			>
avg_disc		Average discount applied per line item.			>
count_order		Number of line items included in the summary gro...			>

**Depends On**

Models

- stg\_tpch\_lineitem

**Code**

# Hosting



## **Azure Blob Storage**

- Store static files in Blob Storage.
- Low cost based on storage and data transfer.

## **Azure Static Web Apps**

- Free tier with 1 custom domain, SSL, and CDN.
- Easy deployment from GitHub/Azure DevOps.
- Ideal for small static sites.

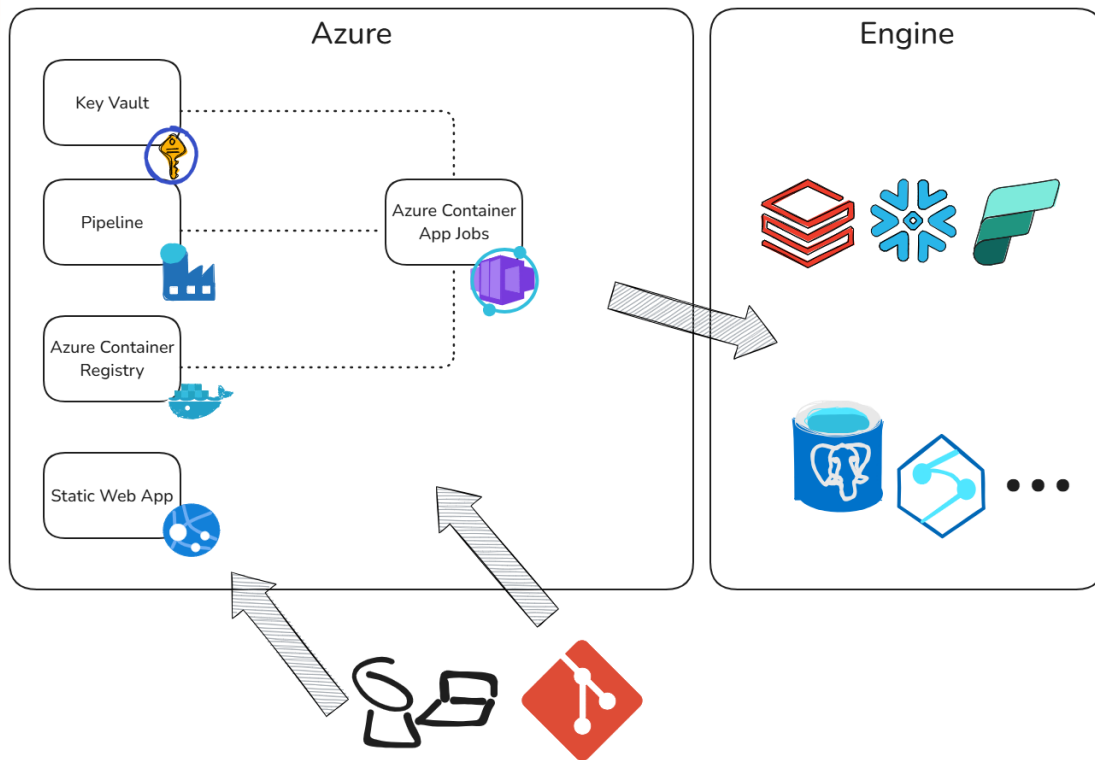
```

35 # install dbt-fabric
36 # -----
37 - script: |
38   | pip install dbt-fabric
39   | displayName: Install dbt-fabric
40
41 # Microsoft ODBC 18
42 # The following sections explain how to install the
43 # Microsoft ODBC driver 18 from the bash shell for different Linux distributions.
44 # https://learn.microsoft.com/en-us/sql/connect/odbc/linux-mac/installing-the-microsoft-odbc-driver-for-sql-server?
45 # -----
46 - task: Bash@3
47   | displayName: Install ODBC 18
48   | inputs:
49   |   targetType: 'inline'
50   |   script: |
51     | curl https://packages.microsoft.com/keys/microsoft.asc | sudo tee /etc/apt/trusted.gpg.d/microsoft.asc
52     | curl https://packages.microsoft.com/keys/microsoft.asc | sudo gpg --dearmor -o /usr/share/keyrings/microsoft-prod.gpg
53     | curl https://packages.microsoft.com/config/ubuntu/\$\(lsb\_release -rs\)/prod.list | sudo tee /etc/apt/sources.list.d/mssql-release.list
54
55     | sudo apt-get update
56     | sudo ACCEPT_EULA=Y apt-get install -y msodbcsql18
57
58 # generate documentation
59 # -----
60 - task: AzureCLI@2
61   | displayName: dbt docs generate
62   | inputs:
63   |   azureSubscription: ${ parameters.serviceconn }
64   |   scriptType: bash
65   |   scriptLocation: inlineScript
66   |   addSpnToEnvironment: true
67   |   workingDirectory: ${ parameters.workingDirectory }
68   |   inlineScript: |
69     | export AZURE_TENANT_ID=$tenantId
70     | export AZURE_CLIENT_ID=$servicePrincipalId
71     | export AZURE_CLIENT_SECRET=$servicePrincipalKey
72
73     | mv profiles.yml.sample profiles.yml
74
75     | dbt deps
76     | dbt docs generate -t ${ parameters.env }

```



# The Poor Man's Setup



# Summary



- If you haven't done so already, take some time to explore concepts like ELT, Analytics Engineering, and the tools that support these roles - such as dbt.
- A production-level setup can be built with dbt Core.
- We can build a generic mechanism that works across different execution engines.
- Databricks makes this much easier for us by providing support for dbt through Databricks Asset Bundles.

# QA & Thanks!